

Data Science

Deriving Knowledge from Data at Scale

Wee Hyong Tok

April 18, 2017



Lecture Outline

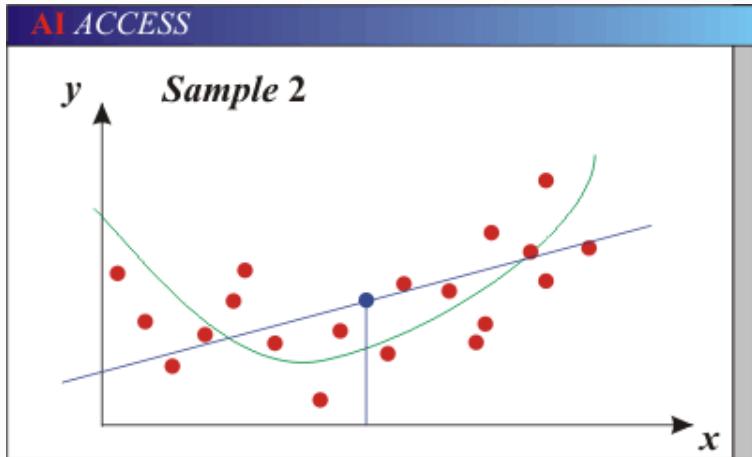
- Assignment 1 Discussion
- Bias and Variance
- Ensembles, Random Forests
- Break
- Lecture 4 Preview
 - Data Science Modelling
 - *Model performance evaluation...*

Continue from last lecture

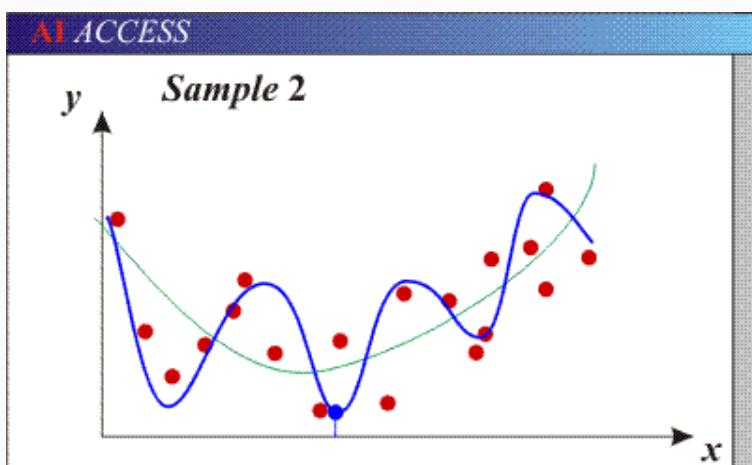
Bias Variance Tradeoff



Bias-Variance Trade-off



- Models with too few parameters are inaccurate because of a large bias (not enough flexibility).



- Models with too many parameters are inaccurate because of a large variance (too much sensitivity to the sample).

Bias-Variance Trade-off

$$E(\text{MSE}) = \text{noise}^2 + \text{bias}^2 + \text{variance}$$

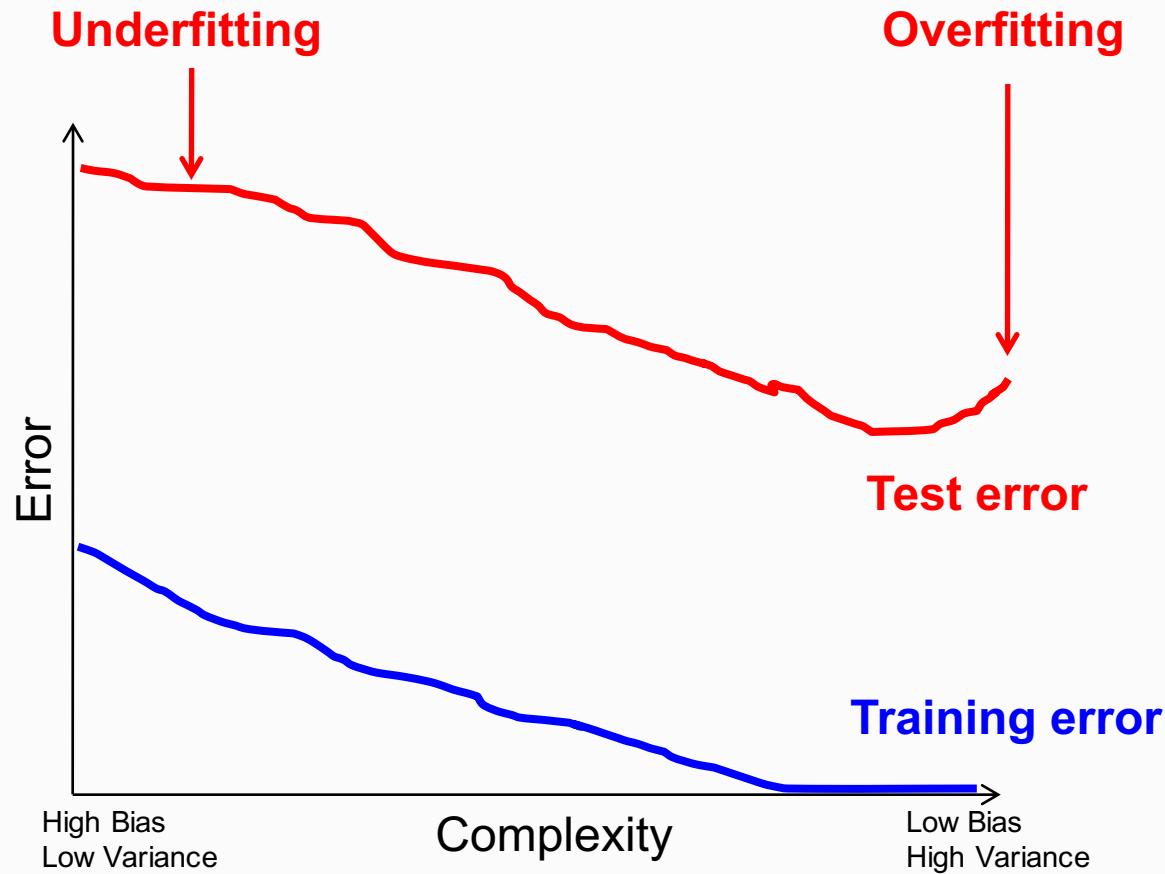
Unavoidable
error

Error due to
incorrect
assumptions

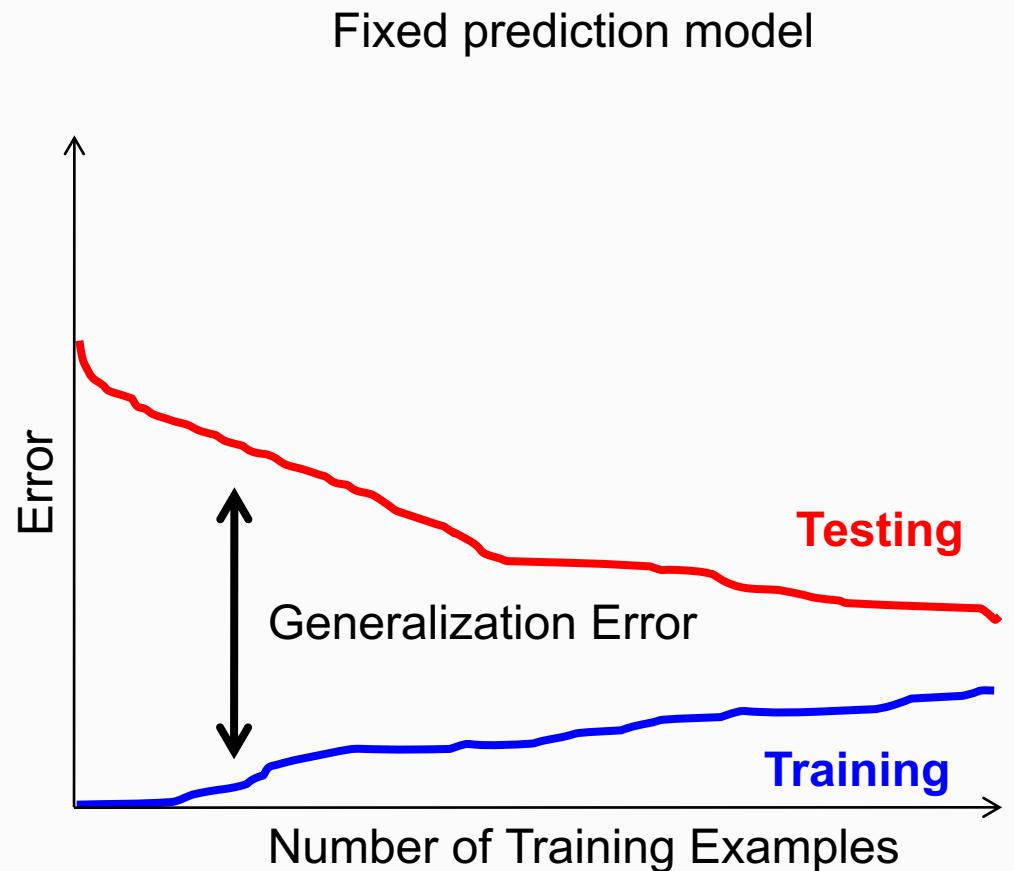
Error due to variance
of training samples

See the following for explanations of bias-variance (also Bishop's "Neural Networks" book):
• <http://www.inf.ed.ac.uk/teaching/courses/mlsc/Notes/Lecture4/BiasVariance.pdf>

Bias-variance tradeoff



Effect of Training Size



Classification



How to reduce variance?

- Choose a simpler classifier
- Regularize the parameters
- Get more training data

What to remember about classifiers

- Try simple classifiers first
- Better to have smart features and simple classifiers than simple features and smart classifiers
- Use increasingly powerful classifiers with more training data (bias-variance tradeoff)

How to be a good data scientist - Be a good modeler...

Example 1: Amazon, big spenders.

Target of the competition was to predict customers who spend a lot of money among customers **using past purchases**.

Data - Transaction data in different categories.

Winning model identified that ‘Free shipping = True’ was an excellent predictor.

Leakage: “*Free Shipping = True*” was simultaneous with the sale, which is a no-no... We can only use data from beforehand to predict the future...

Know your Data

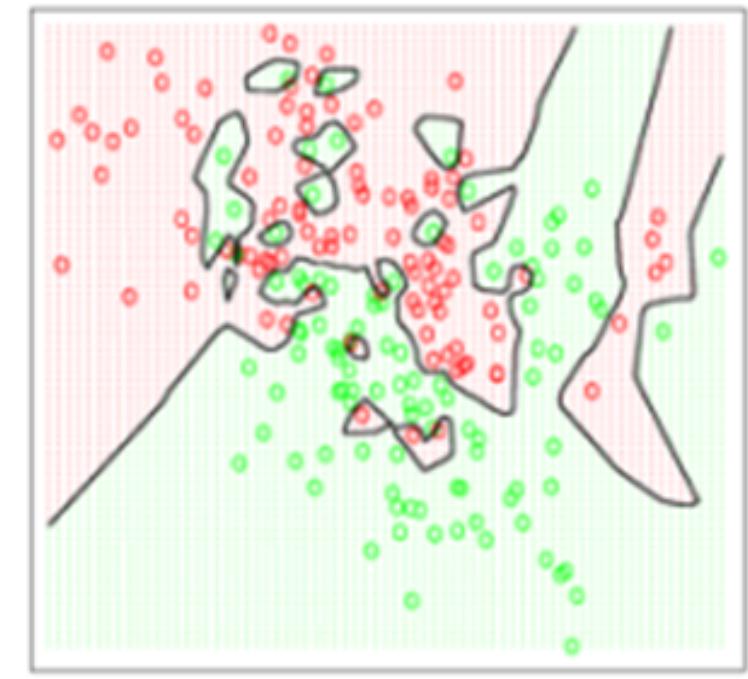
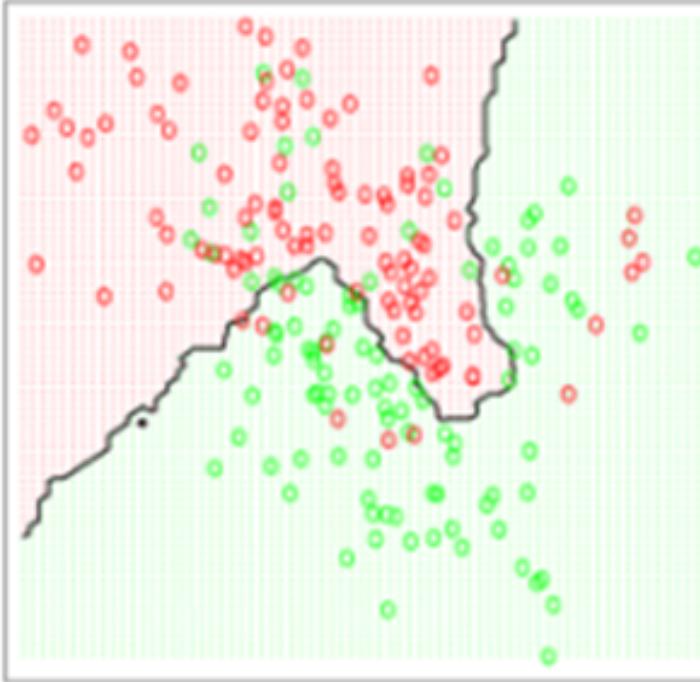
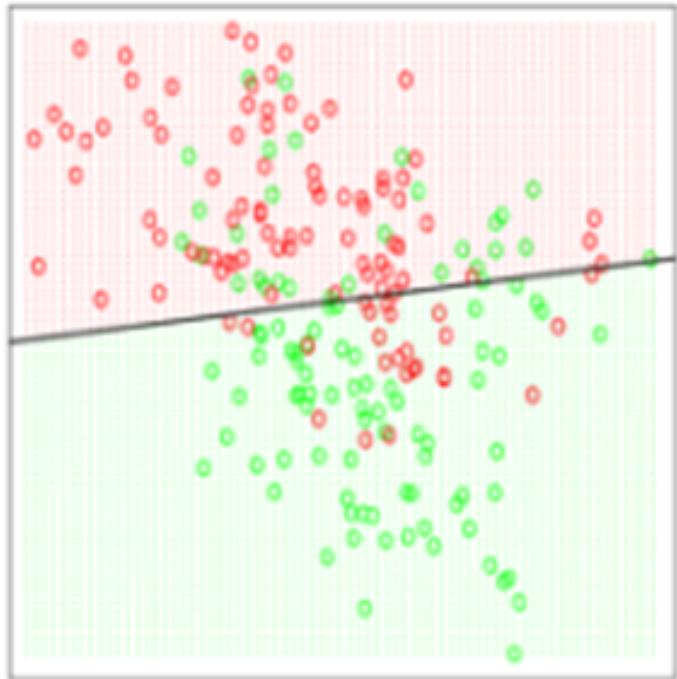
Avoid Leakage

*Winning competition on leakage is easier than building good models. But even if you don't explicitly understand and game the leakage, **your model will do it for you**. Either way, leakage is a huge problem.*

- You need a strict temporal cutoff: remove all information just *prior to the event of interest*.
- There has to be a timestamp on every entry and you need to keep it
- The best practice is to start from scratch with clean, raw data after careful consideration
- You need to know how the data was created! I (try to) work only with data I pulled and prepared myself...

Know Your Model

To avoid overfitting, we cross-validate and we cut down on the complexity of the model to begin with. Here's a standard picture (although keep in mind we generally work in high dimensional space and don't have a pretty picture to look at)



No model is perfect, great models are useful...

The Big Picture

You need to know what the purpose of the model is and how it is going to be used in order to decide how to do it and whether it's actually working...

The art in data science is translating the problem into the language of data, features, proxy variable(s), a prediction,...

The science in data science is given raw data, constraints and a problem statement, you have an infinite set of models to choose from, with which you can use to maximize performance on ***some evaluation metric***. Every design choice you make can be formulated as a hypothesis, upon which you will use ***rigorous testing and experimentation to either validate or refute***.

The Big Picture

Given

- data
- a problem, and
- constraints

We need to determine

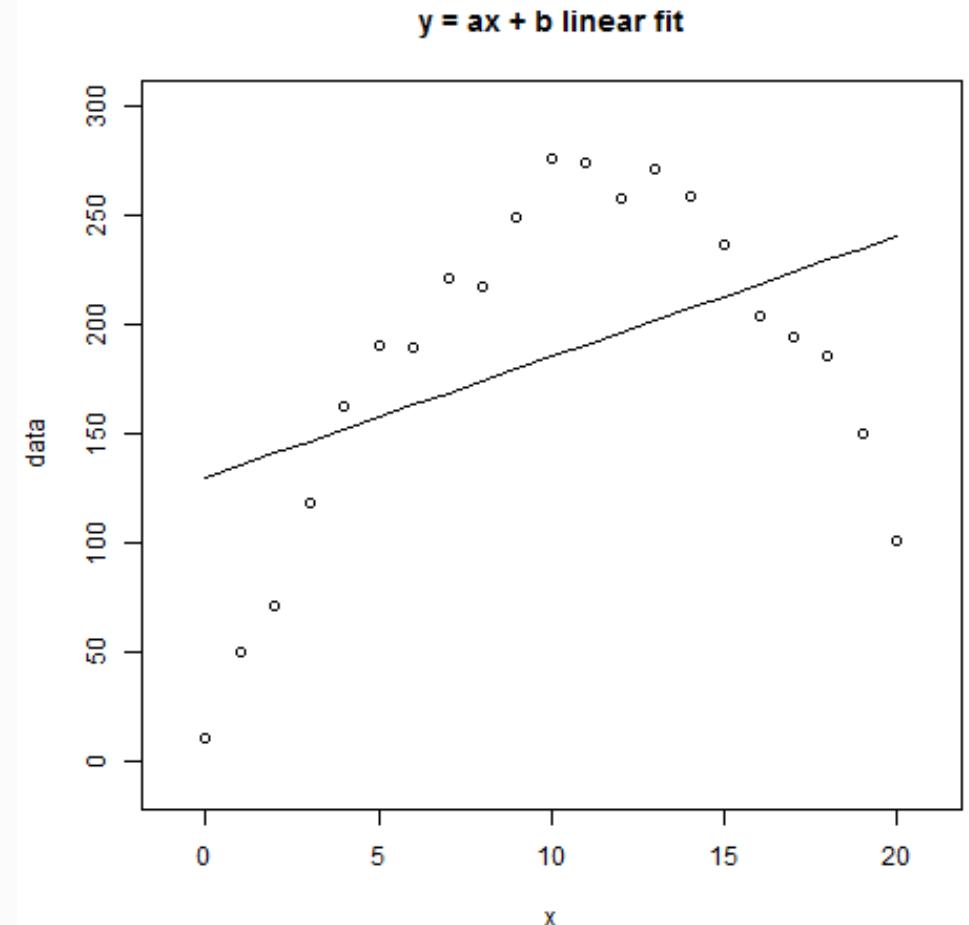
- features,
- a classifier,
- an optimization method, and
- an evaluation metric.

Later today we will focus on **evaluation metrics...**



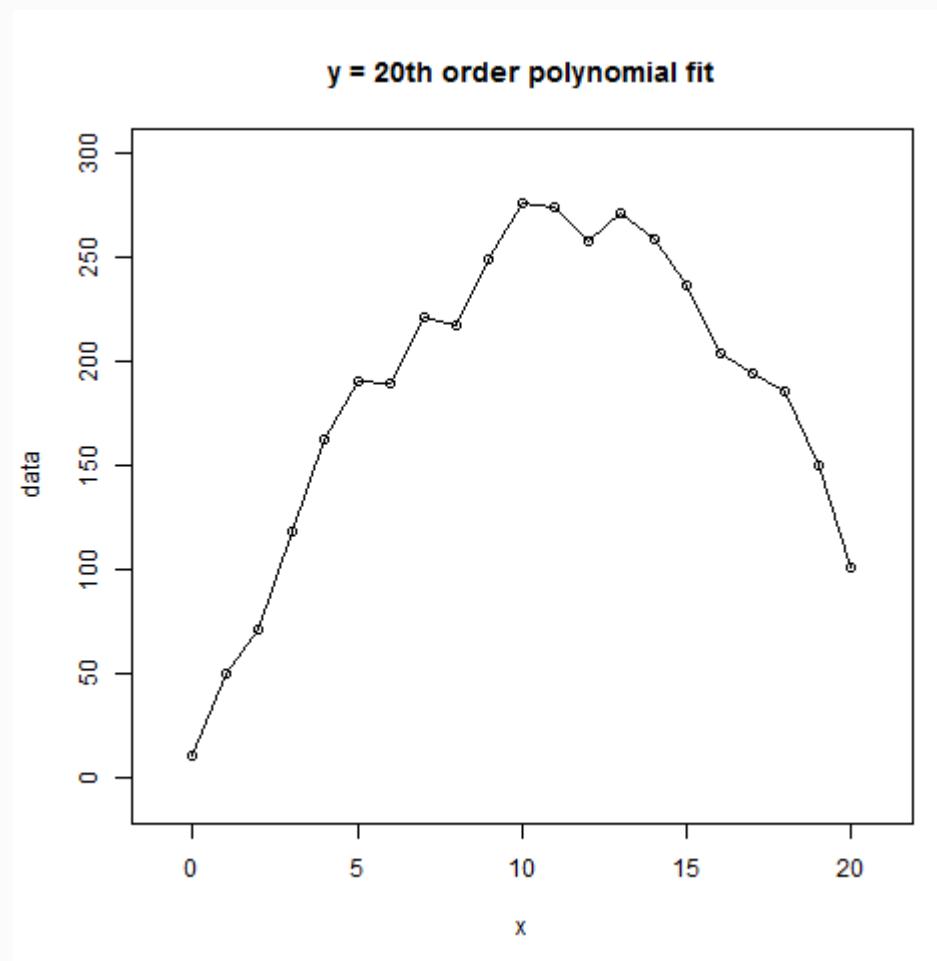
Review: Under-Fitting

- Under fit model to the data
 - Model is not expressive enough to capture the (quadratic polynomial) signal in the data
 - Big approximation errors in training data (high bias)
 - Little variance between multiple test datasets during prediction (low variance)



Review: Over-Fitting

- Over fit model to the data
 - Model has just learnt every point in the training data
 - No approximation errors on training data (low bias)
 - Generalize poorly across multiple test datasets during prediction (high variance)

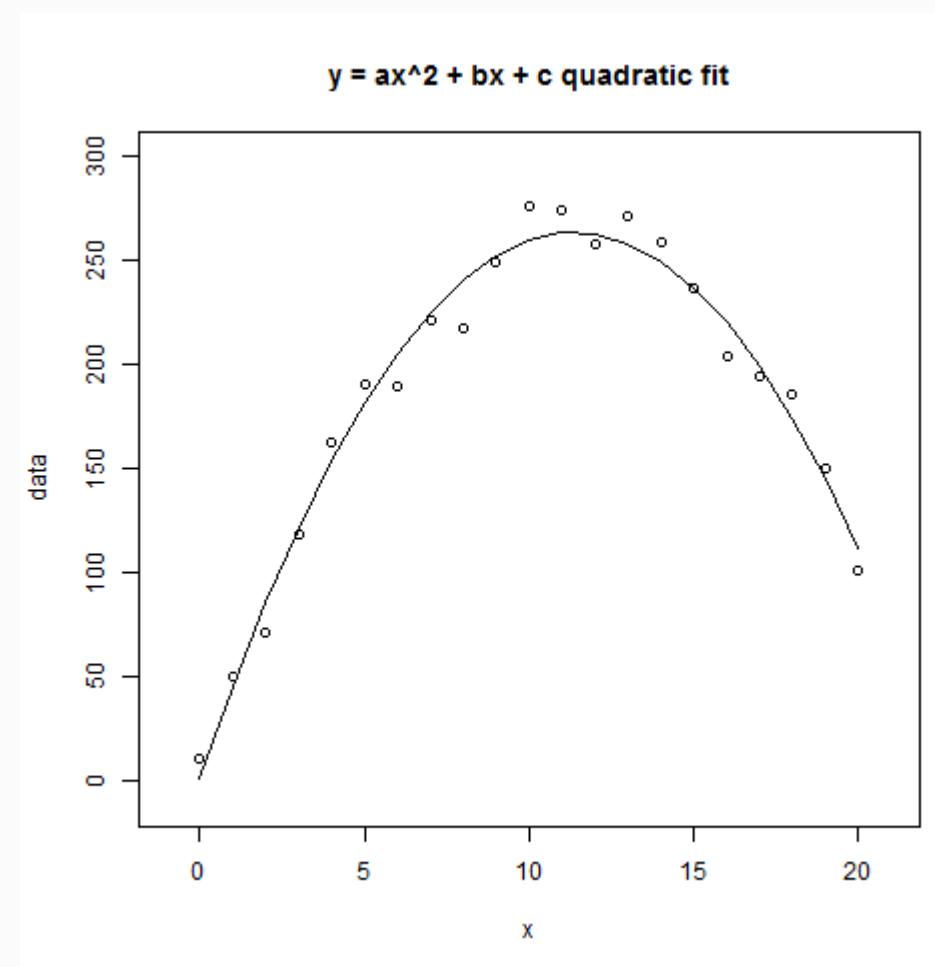


Bias-Variance Trade-Off

Trade-off between

- Capturing high level structure in data
 - Reduce bias
 - Provides some variation in new datasets
- Not memorize the training data
 - Reduce variance
 - Generalize to new datasets

Trade-off is called “Bias-Variance” trade-off



Reducing Under Fitting

- Increase model complexity, for e.g.
 - Increase the number of levels in a decision tree
 - Increase the number of hidden layers in a neural network.
 - Decrease the number of neighbors (k) in k-NN
- Increase the number of features
- In iterative training algorithms, iterate long enough so that the objective function has converged.

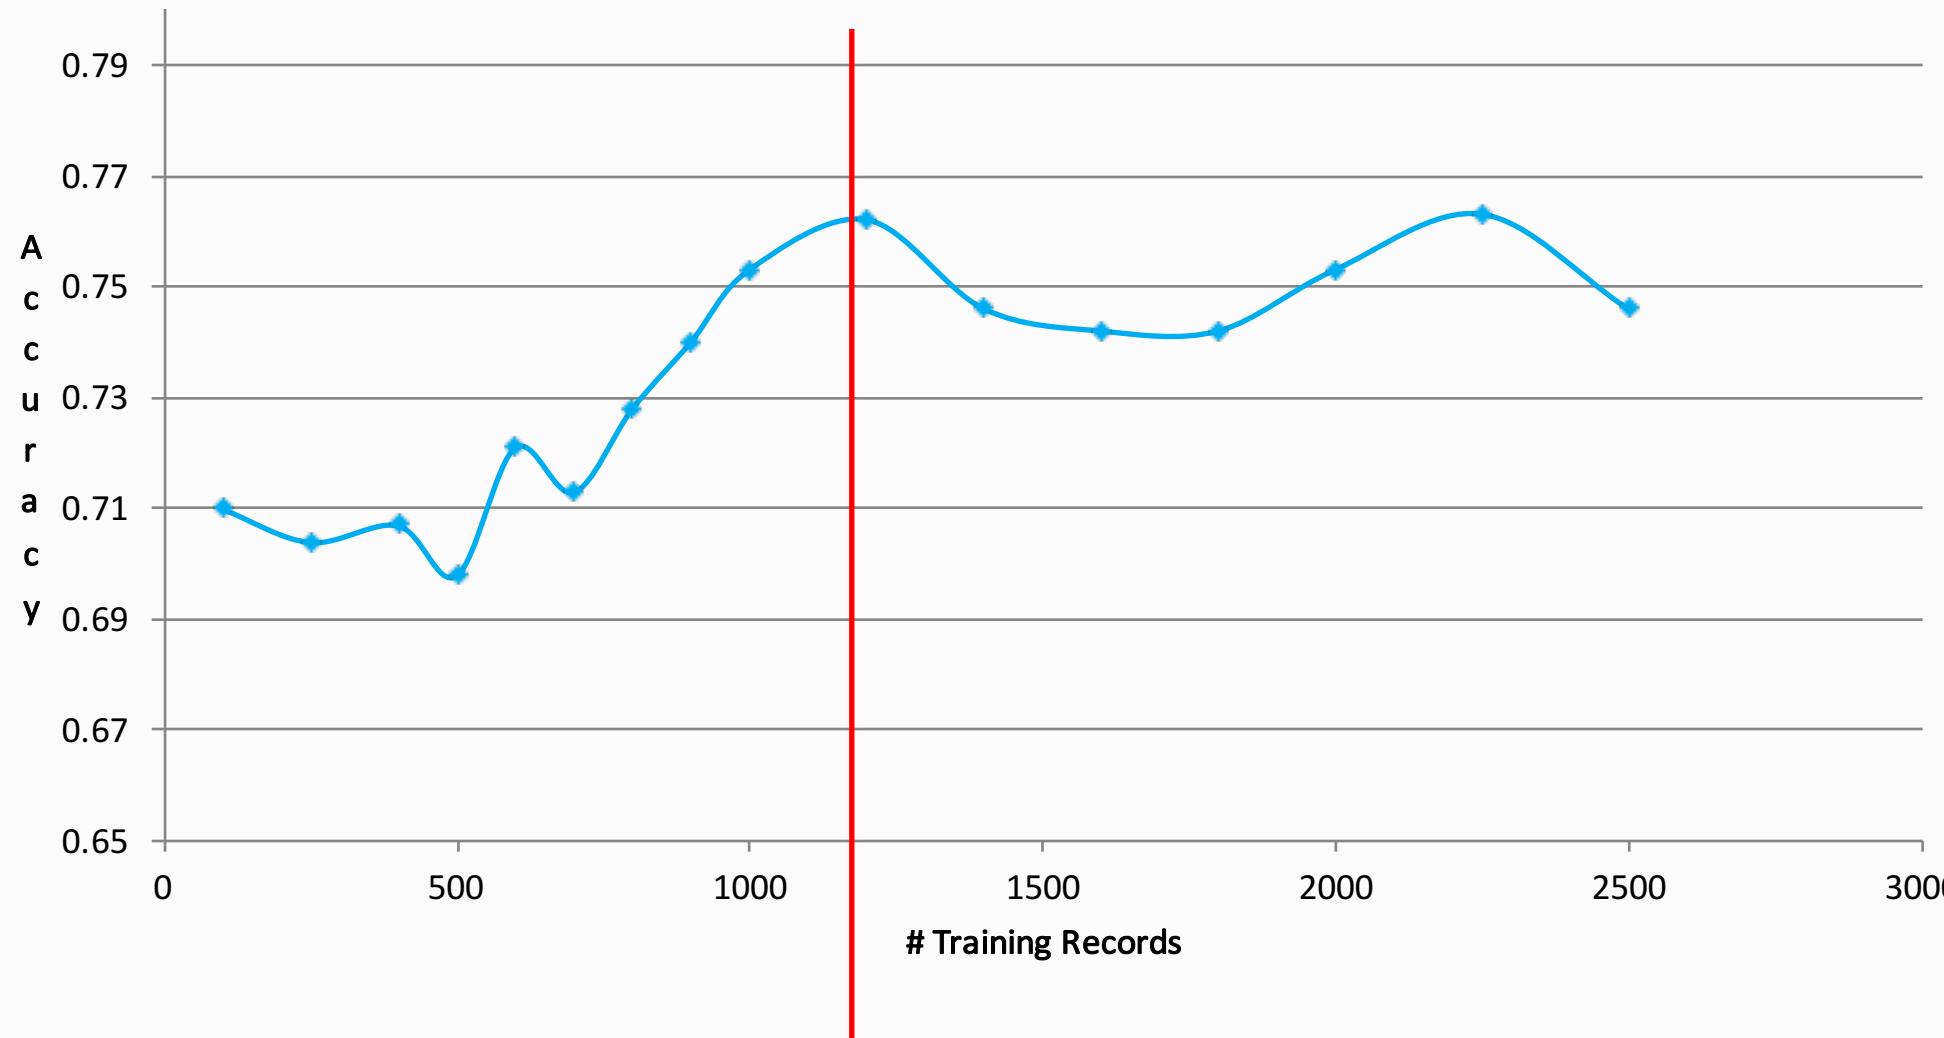


Reducing Over Fitting

- Decrease model complexity, for e.g.
 - Prune a decision tree
 - Reduce the number of hidden layers in a neural network.
 - Increase the number of neighbors (k) in k-NN
- Decrease the number of features
 - More aggressive feature selection
- Regularization (control feature complexity)
 - Penalize high weights.
 - L-1 regularization (LASSO) very efficient at pushing weights of non-informative features to 0.
- Gather more training data if possible
- In iterative training algorithms, stop training earlier to prevent “memorization” of training data



Sufficiency of Training Data (Example)



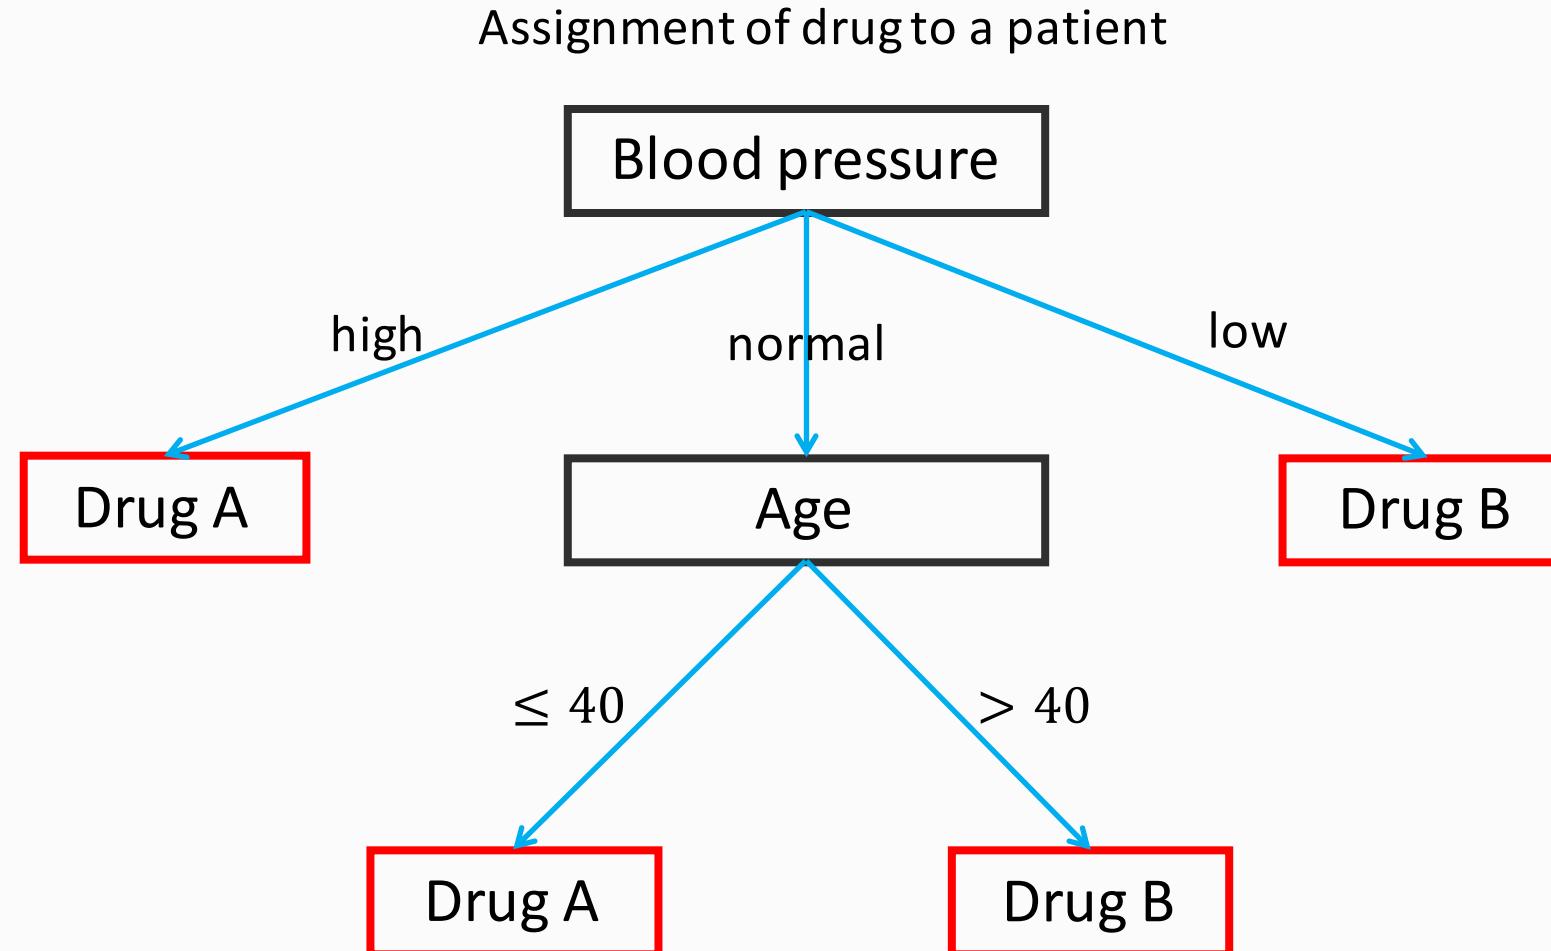
Accuracy on test data stabilizes above 1000 training samples

Review: Decision Tree

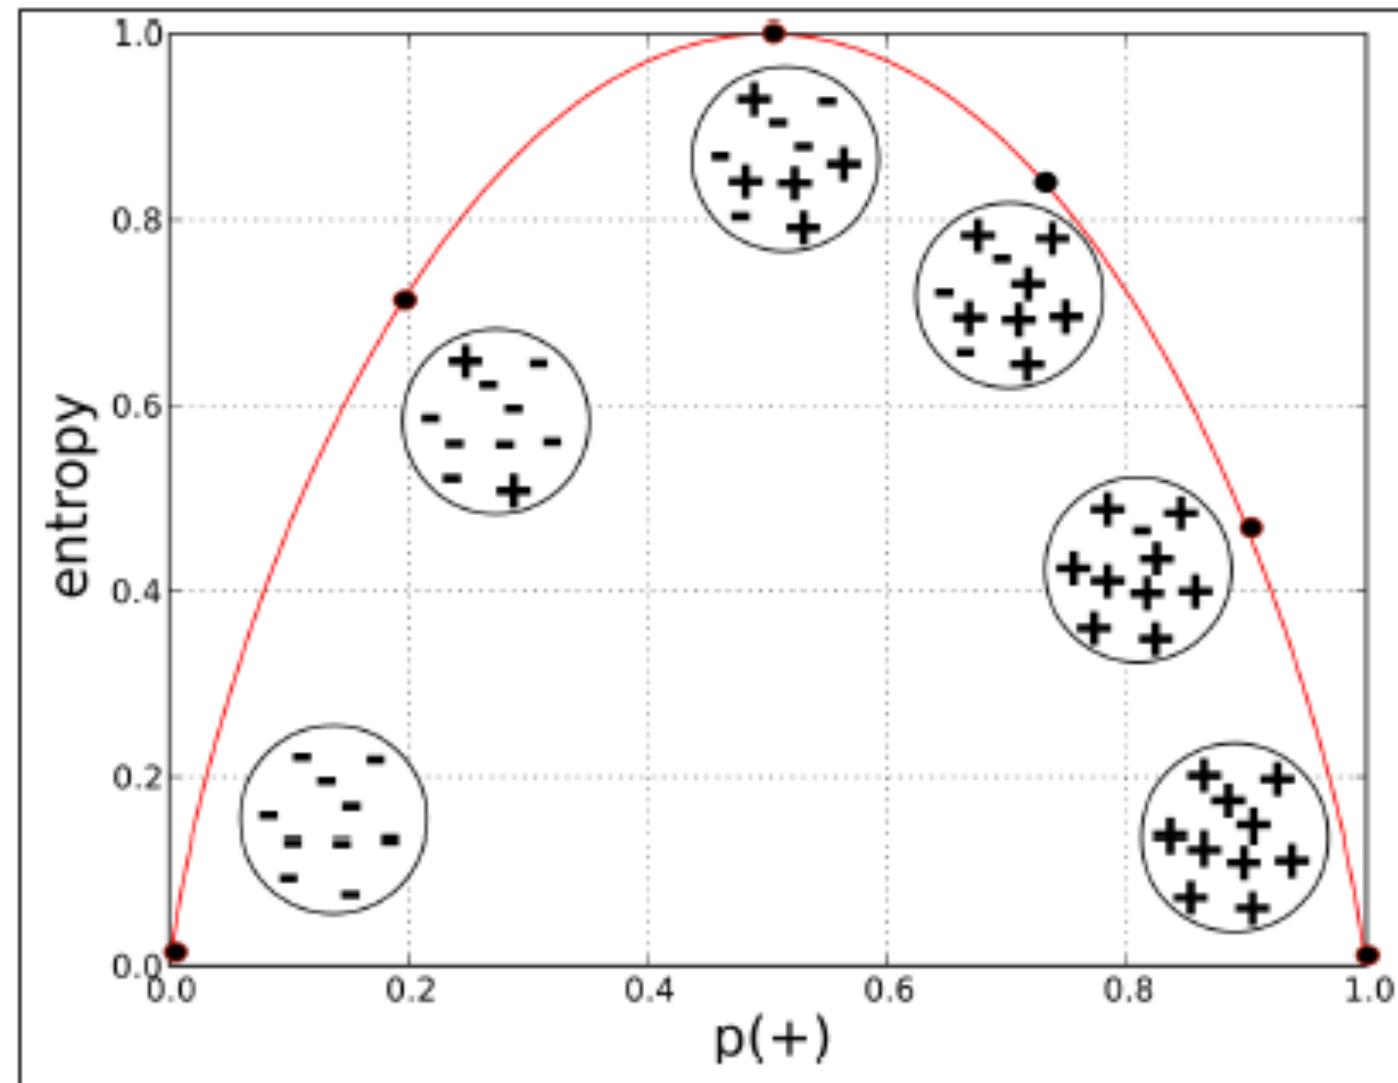
Unique Features

1. Automatically selects features
2. Able to handle large number of features
3. Numeric, nominal, missing
4. Easy to ensemble (Random Forrest, Boosted DT)
5. Transparent and easily explainable😊...

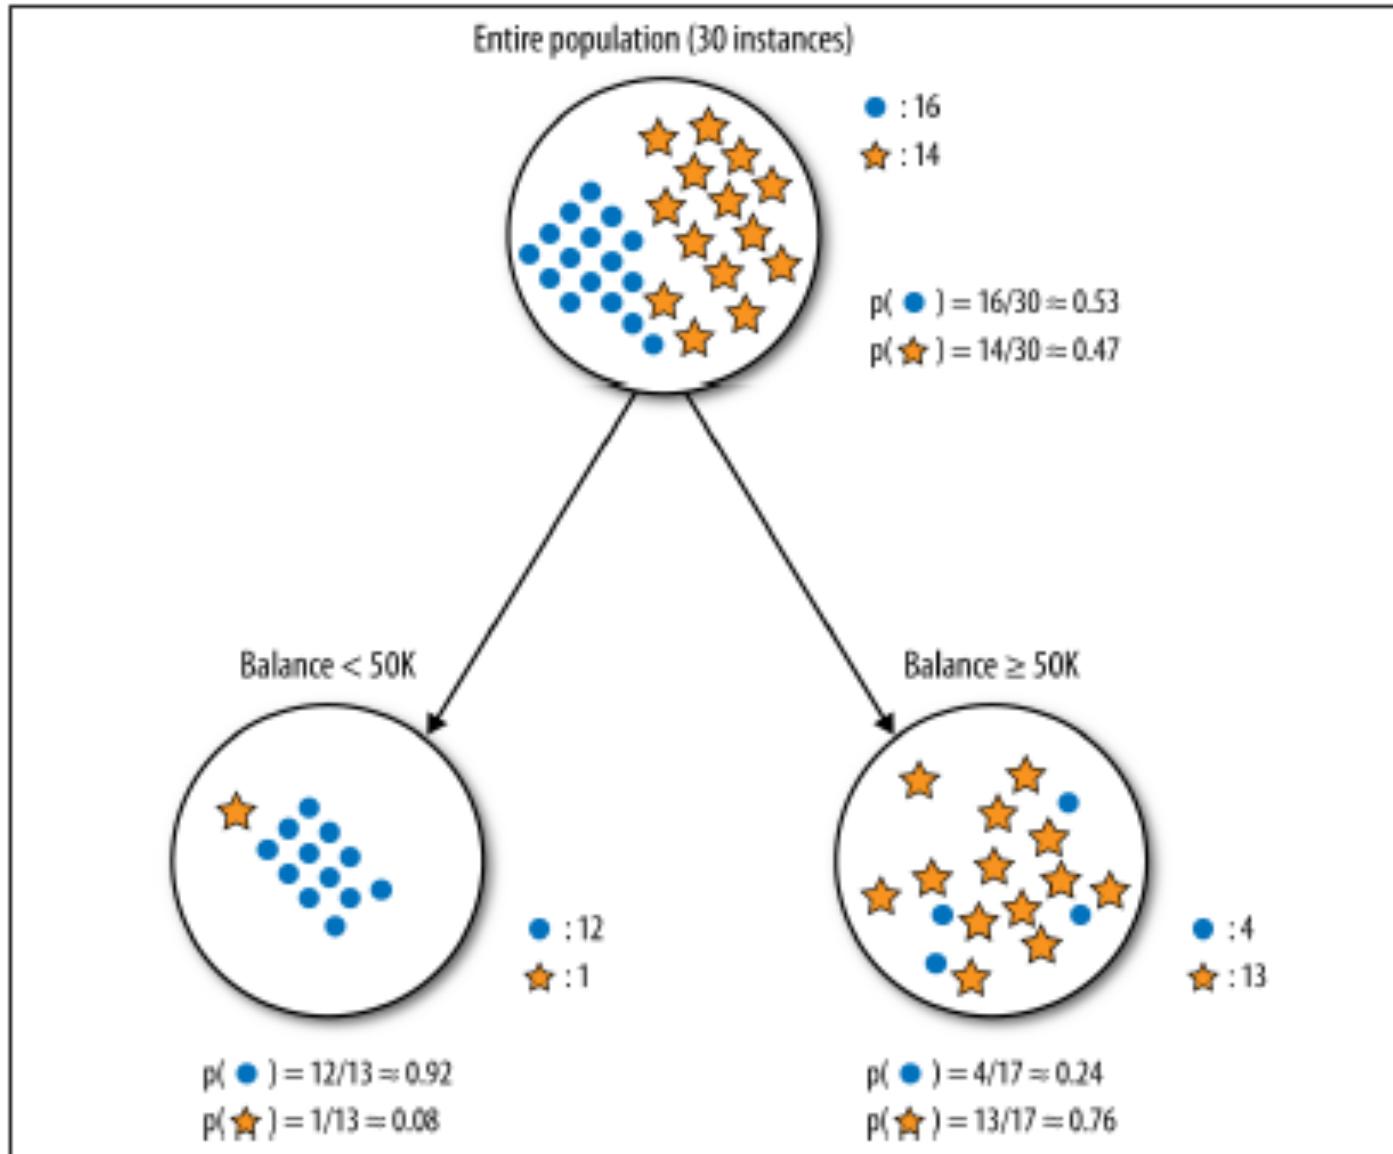
Review: Decision Tree



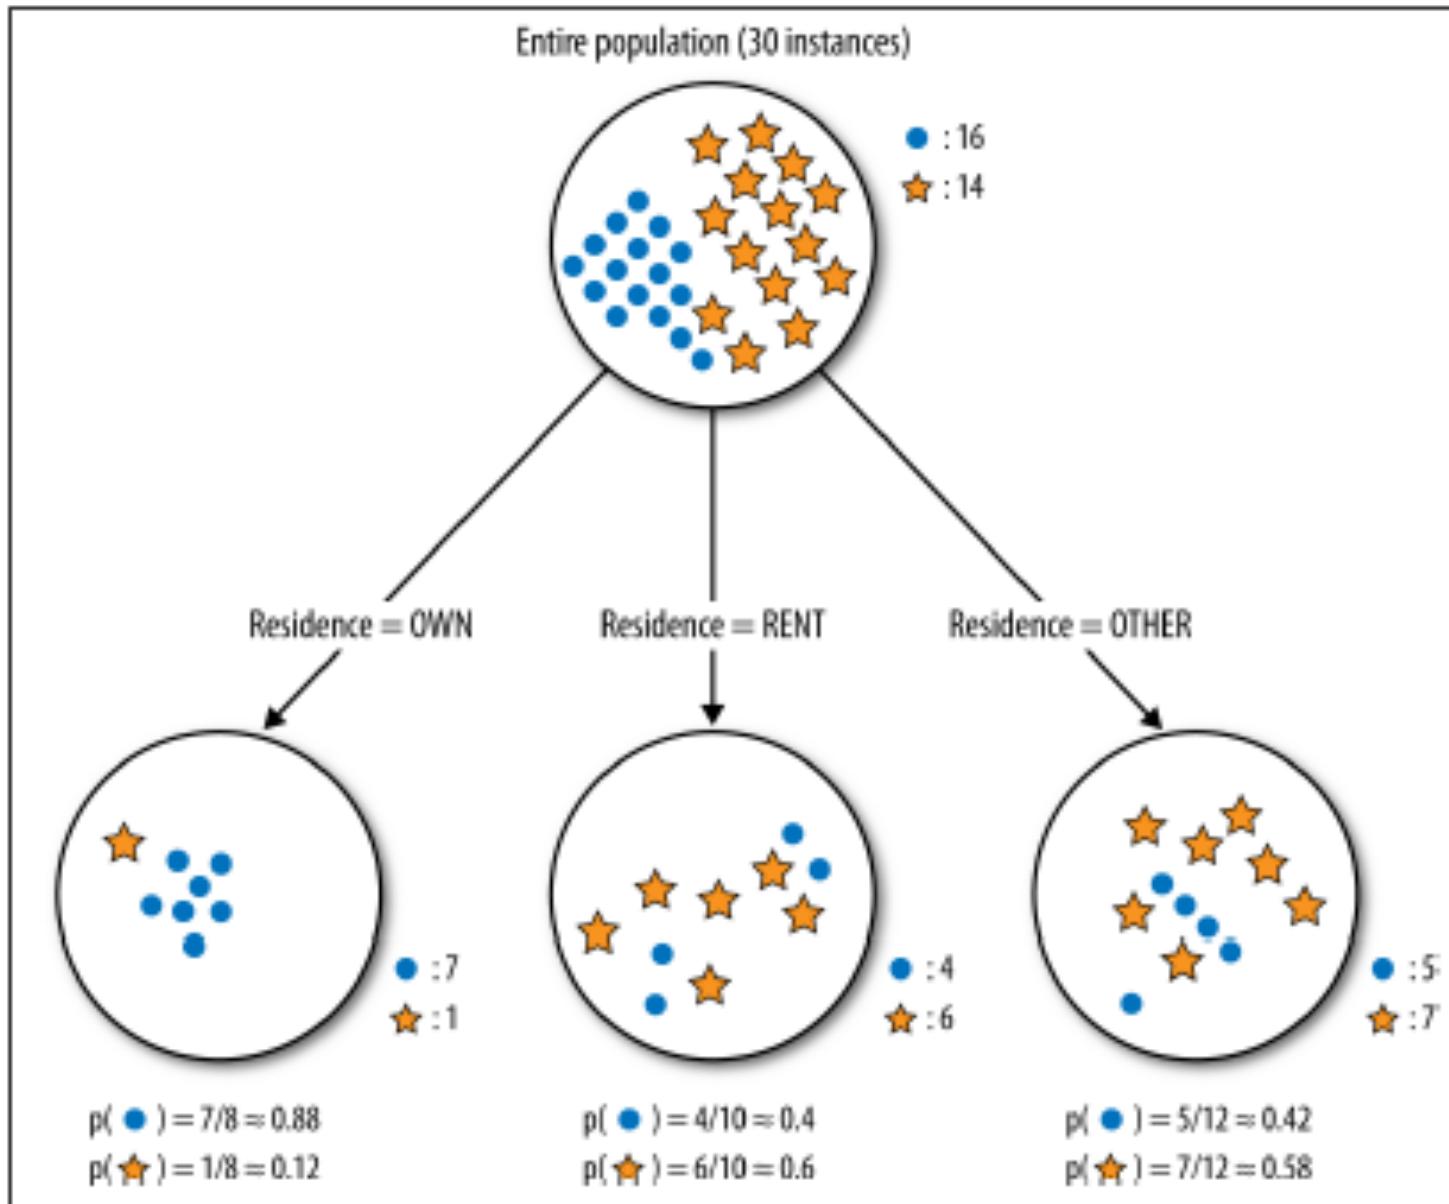
It's all about minimizing the entropy (variance)...



It's all about minimizing the entropy (variance)...



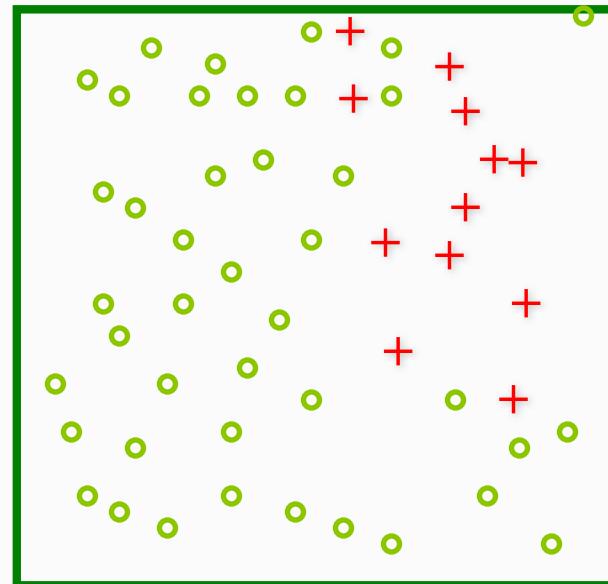
It's all about minimizing the entropy (variance)...



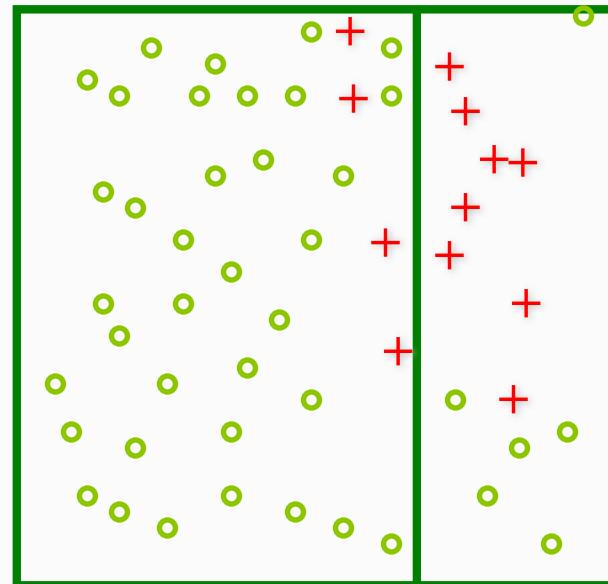
Review: Induction of decision trees

- Top down approach
 - Build the decision tree from top to bottom, from the root to the leaves
- Greedy selection of a test feature
 - Compute an evaluation measure for all features
 - Select the feature with the best measure
- Divide and Conquer/ Recursive Descent
 - Divide examples according to values of test feature
 - Apply the procedure recursively to the subsets
 - Terminate the recursion if
 - All cases belong to the same class, no more examples are available, cutoff condition has been satisfied (minimum node size)

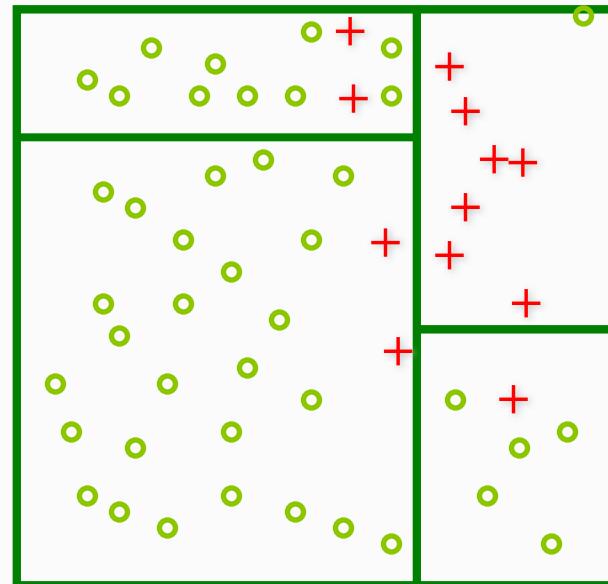
Review: Spatial example, recursive binary splits



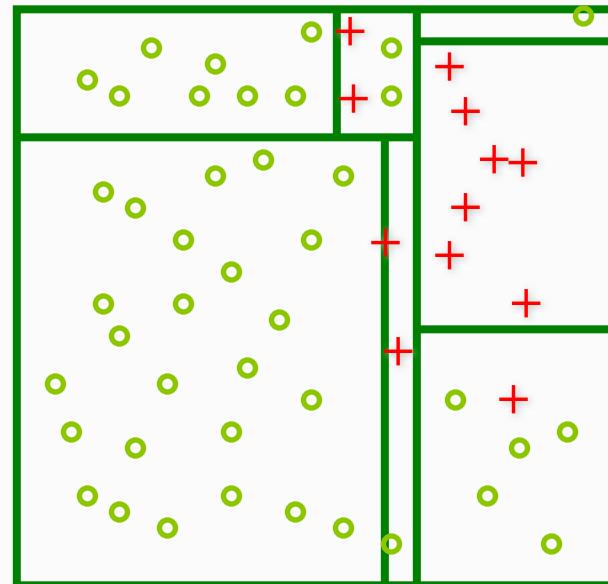
Review: Spatial example, recursive binary splits



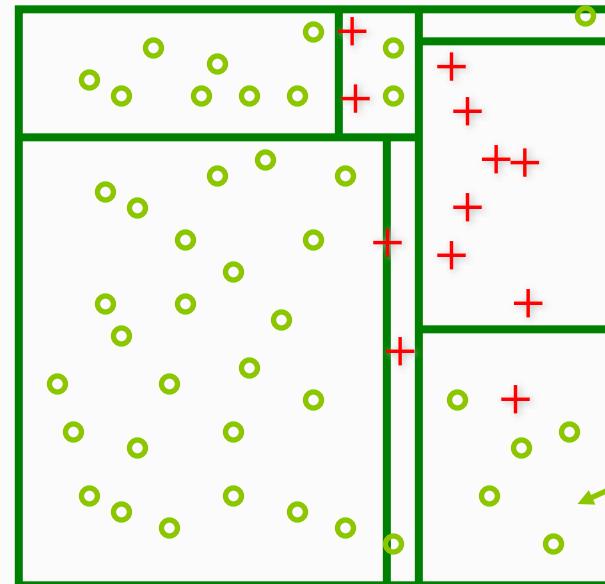
Review: Spatial example, recursive binary splits



Review: Spatial example, recursive binary splits



Review: Spatial example, recursive binary splits



Once regions are chosen class probabilities are easy to calculate

$$p_m = 5/6$$

Ensemble and Random Forests



Random Forest (Decision Forests)

Learning ensemble consisting of a **bagging** of un-pruned decision tree learners with a randomized selection of features at each split.

Decision trees are the individual learners that are combined. Decision trees are one of most popular learning methods commonly used for data exploration

- ❖ Powerful, explainable (decision trees, that is...)
- ❖ Automatic Feature Selection
- ❖ Ensembling



Why Ensemble with Decision Trees?

Decision trees: Advantages and limitations

Advantages:

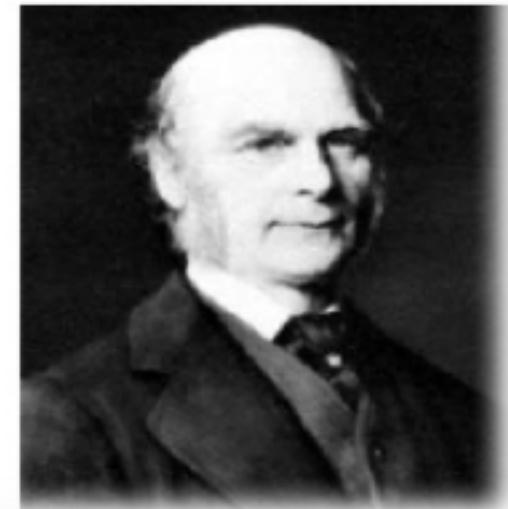
Characteristic	Neural nets	SVM	Trees	MARS	k-NN, kernels
Natural handling of data of “mixed” type	●	●	●	●	●
Handling of missing values	●	●	●	●	●
Robustness to outliers in input space	●	●	●	●	●
Insensitive to monotone transformations of inputs	●	●	●	●	●
Computational scalability (large N)	●	●	●	●	●
Ability to deal with irrelevant inputs	●	●	●	●	●
Ability to extract linear combinations of features	●	●	●	●	●
Interpretability	●	●	●	●	●
Predictive power	●	●	●	●	●

Limitations:

- Low prediction accuracy
- High variance
- Ensemble, to maintain advantages while increasing accuracy !

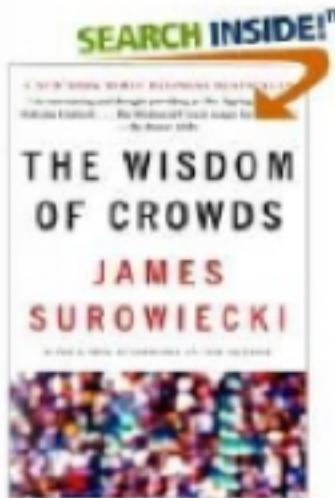
Francis Galton

- Galton promoted statistics and invented the concept of correlation.
- In 1906 Galton visited a livestock fair and stumbled upon an intriguing contest.
- An ox was on display, and the villagers were invited to guess the animal's weight.
- Nearly 800 gave it a go and, not surprisingly, not one hit the exact mark: 1,198 pounds.
- Astonishingly, however, the average of those 800 guesses came close - very close indeed. It was 1,197 pounds.



The Wisdom of Crowds

Why the Many Are Smarter Than the Few and How Collective Wisdom Shapes Business, Economies, Societies and Nations



- Under certain controlled conditions, the aggregation of information in groups, resulting in decisions that are often superior to those that can be made by any single - even experts.
- Imitates our second nature to seek several opinions before making any crucial decision. We weigh the individual opinions, and combine them to reach a final decision

Does it work...

Not all crowds (groups) are wise...

Example: investors in a stock market bubble (swing)



Key Criteria

Diversity of Opinion

- Each person should have private information even if it's just an eccentric interpretation of the known facts.

Independence

- People's opinions are not determined by the opinions of those around them.

Decentralization

- People are able to specialize and draw on local knowledge.

Aggregation

- Some mechanism exists to turn private judgments into a collective decision

Random forests

- Random forests (RF) are a combination of decision tree predictors
- Each tree depends on the values of a random vector sampled independently and with the same distribution for all trees in the forest
- The generalization error of a forest of tree classifiers depends on the strength of the individual trees in the forest and the correlation between them
- Using a random selection of features to split each node yields error rates that compare favorably to Adaboost, and are more robust with respect to noise

Random forests

D = training set

k = nb of trees in forest

F = set of tests

n = nb of tests

for $i = 1$ to k do:

 build data set D_i by sampling with replacement from D

 learn tree T_i (\tilde{T}_i) from D_i :

 at each node:

 choose best split from random subset of F of size n

 allow aggregates and refinement of aggregates in tests

make predictions according to majority vote of the set of k trees.

Random Forests

- Ensemble method tailored for decision tree classifiers
- Creates k decision trees, where each tree is independently generated based on random decisions
- Bagging using decision trees can be seen as a special case of random forests where the random decisions are the random creations of the bootstrap samples

Two examples of random decisions in RFs

- At each internal tree node, randomly select F attributes, and evaluate just those attributes to choose the partitioning attribute
 - Tends to produce trees larger than trees where all attributes are considered for selection at each node, but different classes will be eventually assigned to different leaf nodes, anyway
 - Saves processing time in the construction of each individual tree, since just a subset of attributes is considered at each internal node
- At each internal tree node, evaluate the quality of all possible partitioning attributes, but randomly select one of the F best attributes to label that node (based on InfoGain, etc.)
 - Unlike the previous approach, does not save processing time



Random forest: first randomization through bagging

Bagging or *bootstrap aggregation* is a technique for reducing the variance of an estimated prediction function.

Bootstrap sample = create new training sets by random sampling the given one $N' \leq N$ times with replacement

Bootstrap aggregation, parallel combination of learners, independently trained on distinct bootstrap samples

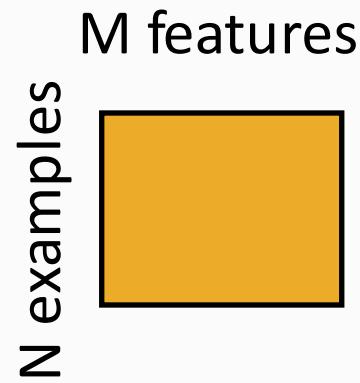
Final prediction is the mean prediction (regression) or class with maximum votes (classification).

Using squared error loss, bagging alone decreases test error by lowering prediction variance, while leaving bias unchanged.

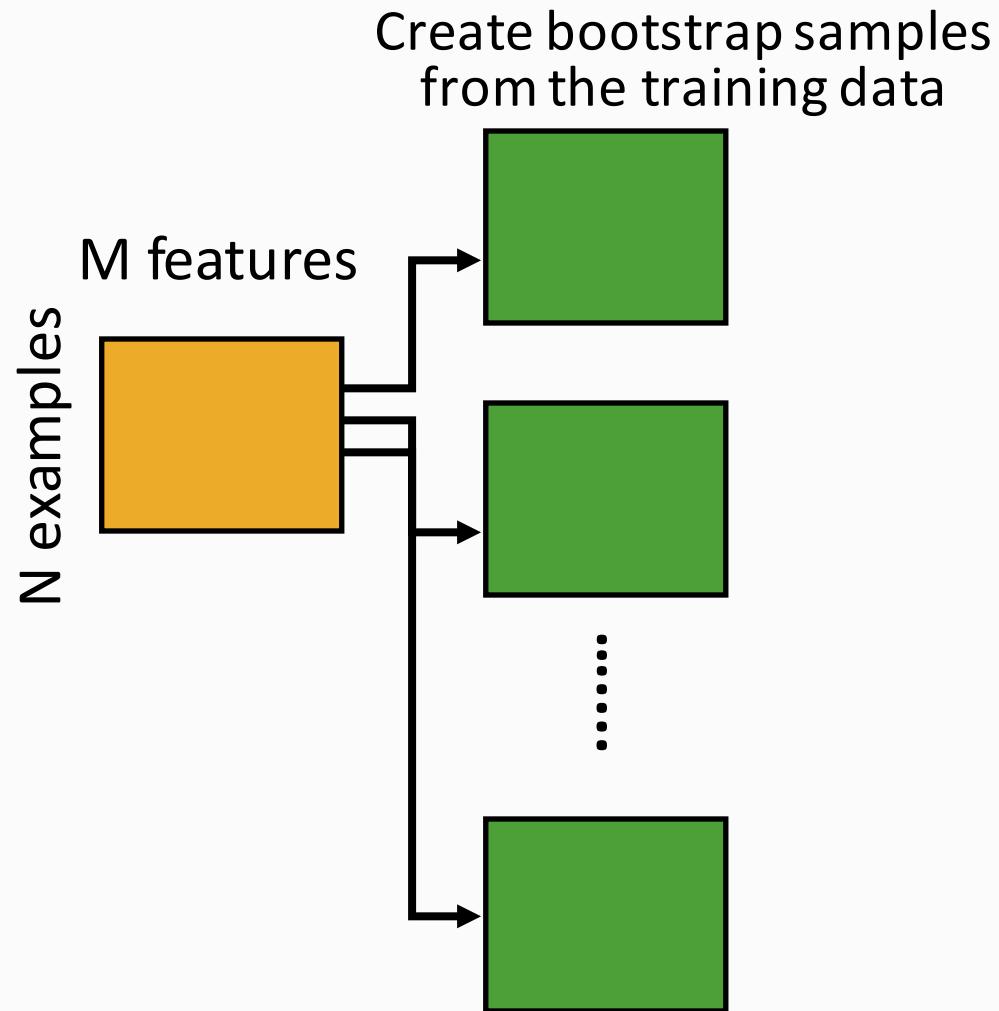


Random Forest Classifier

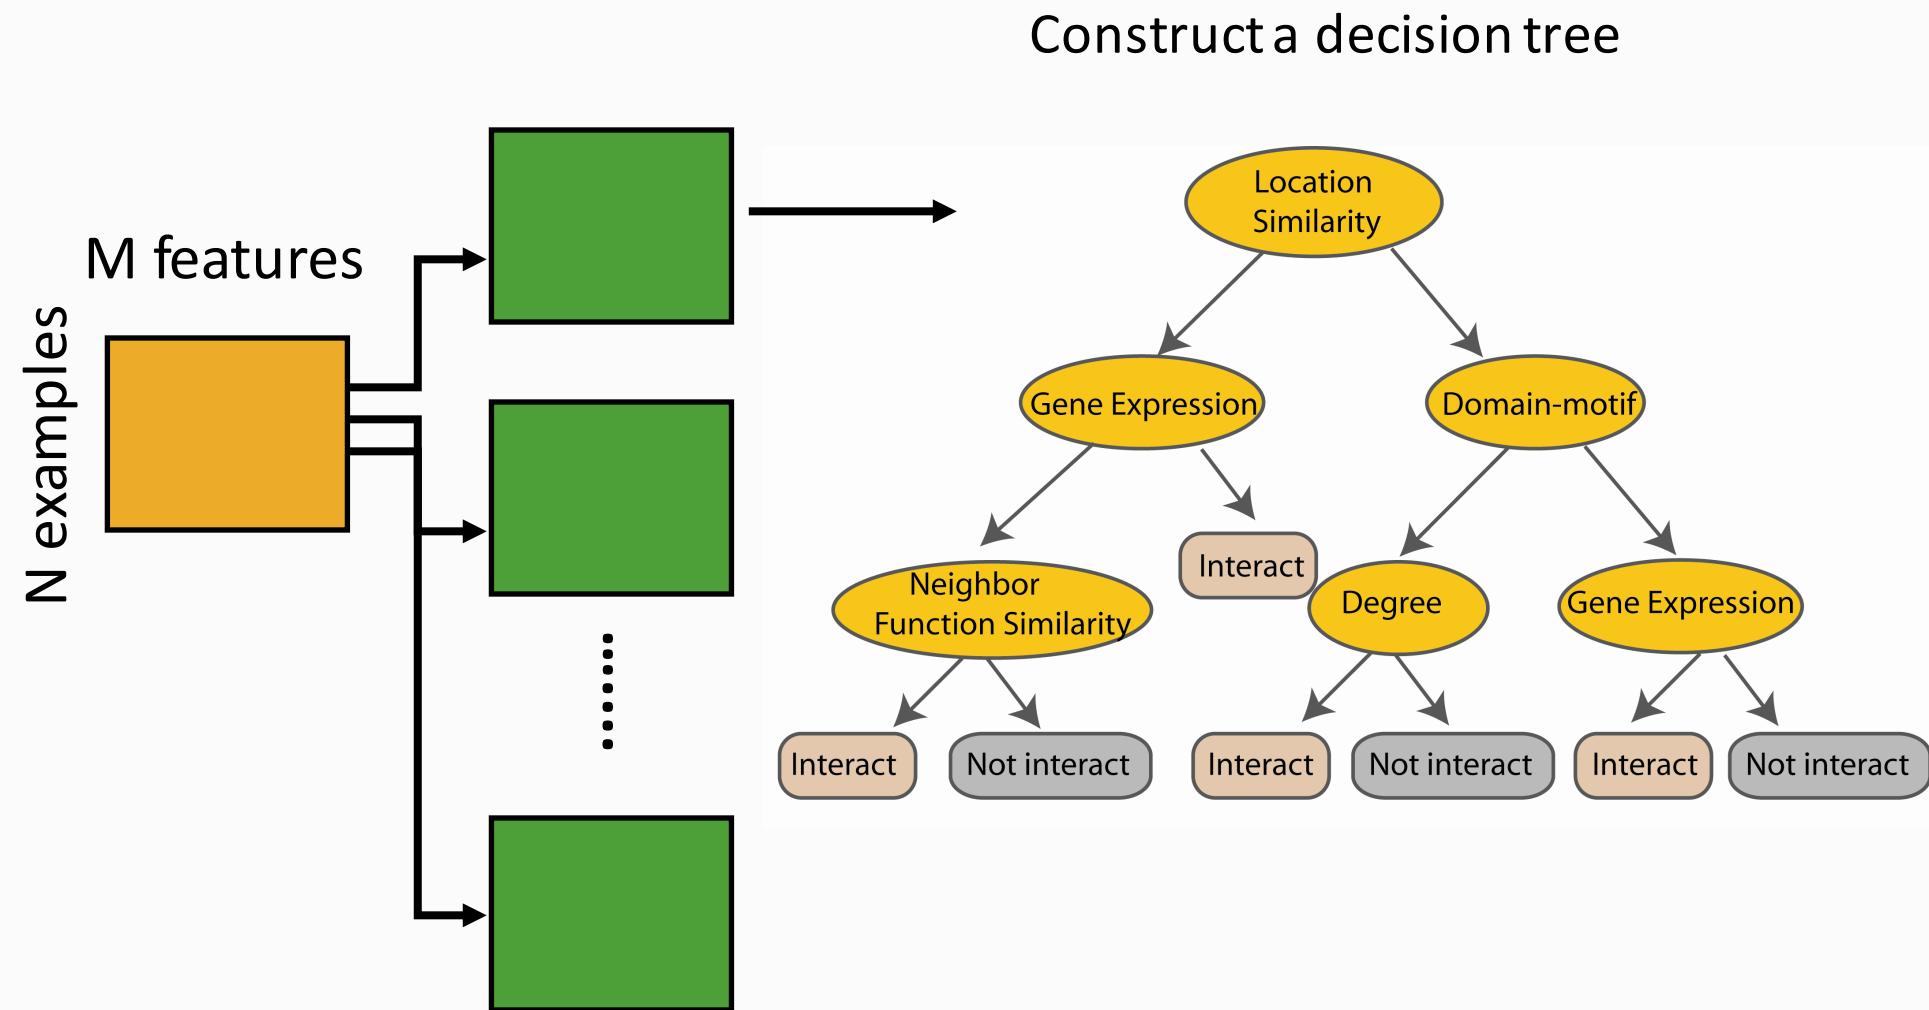
Training Data



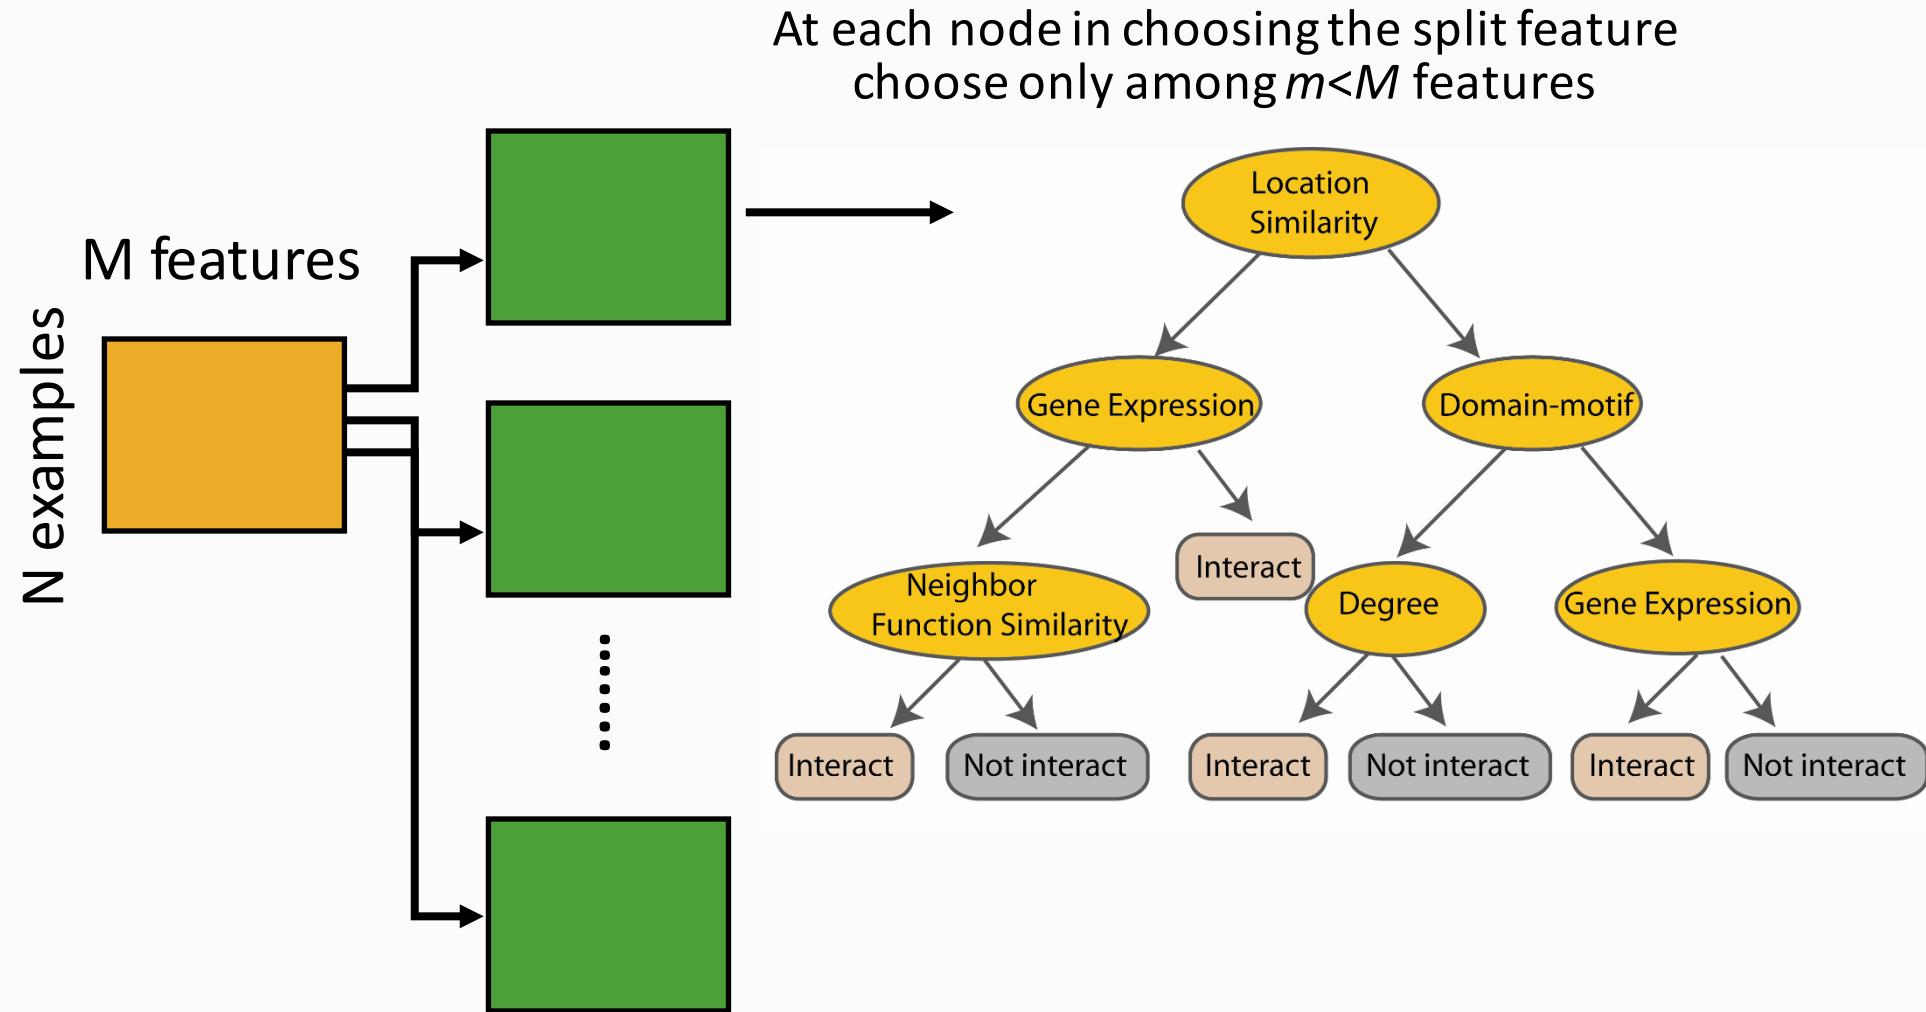
Random Forest Classifier



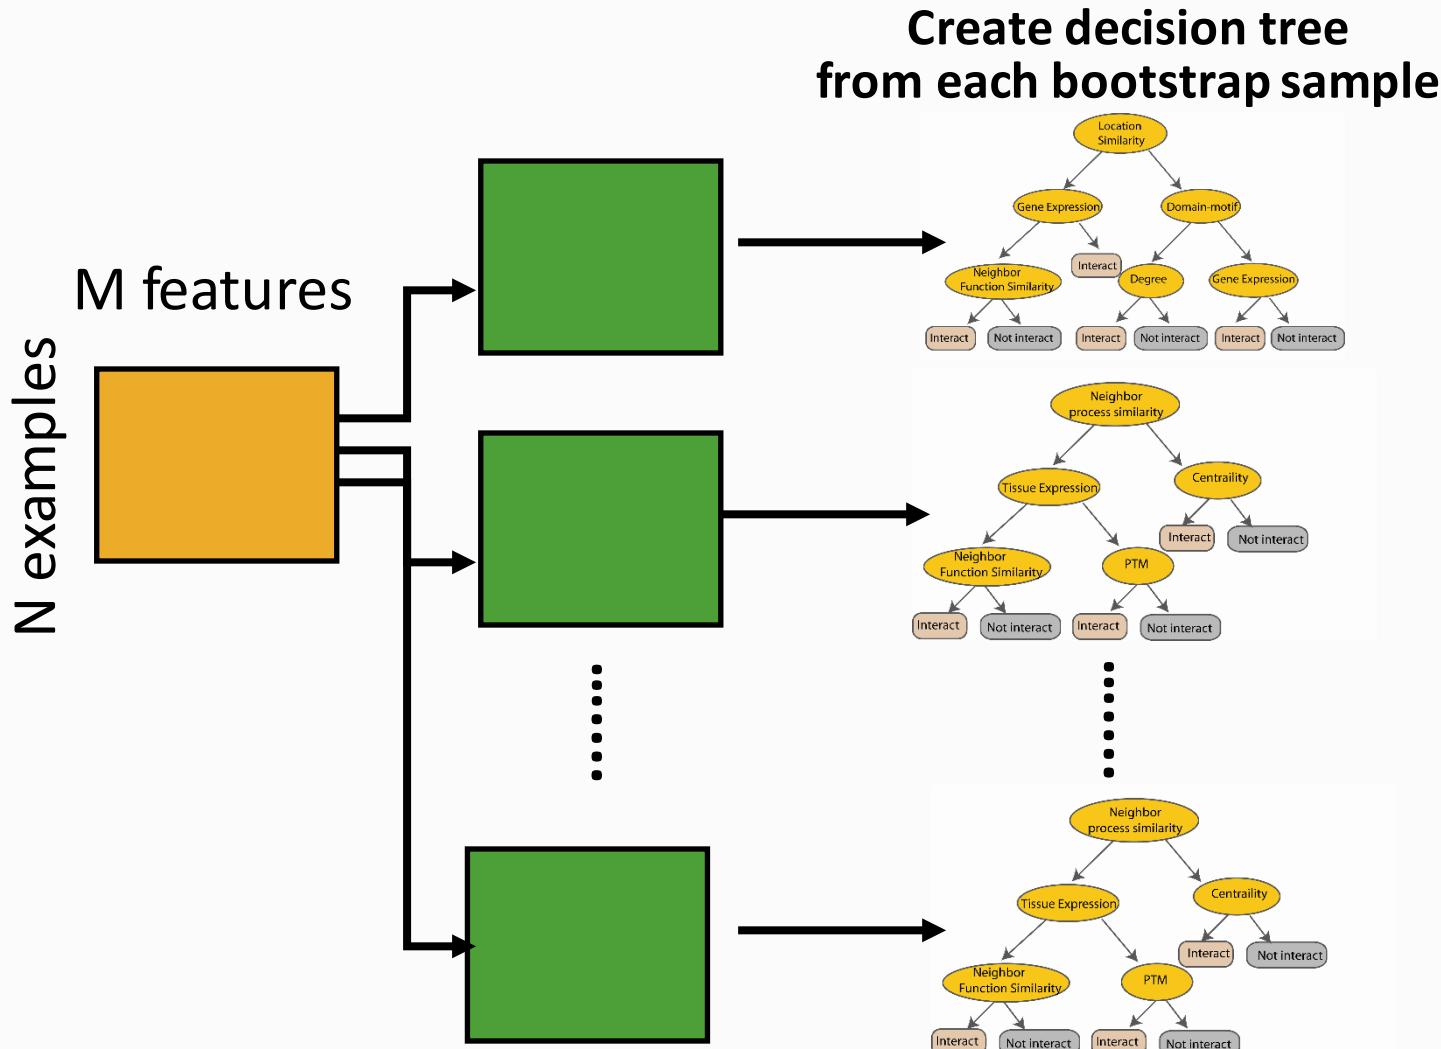
Random Forest Classifier



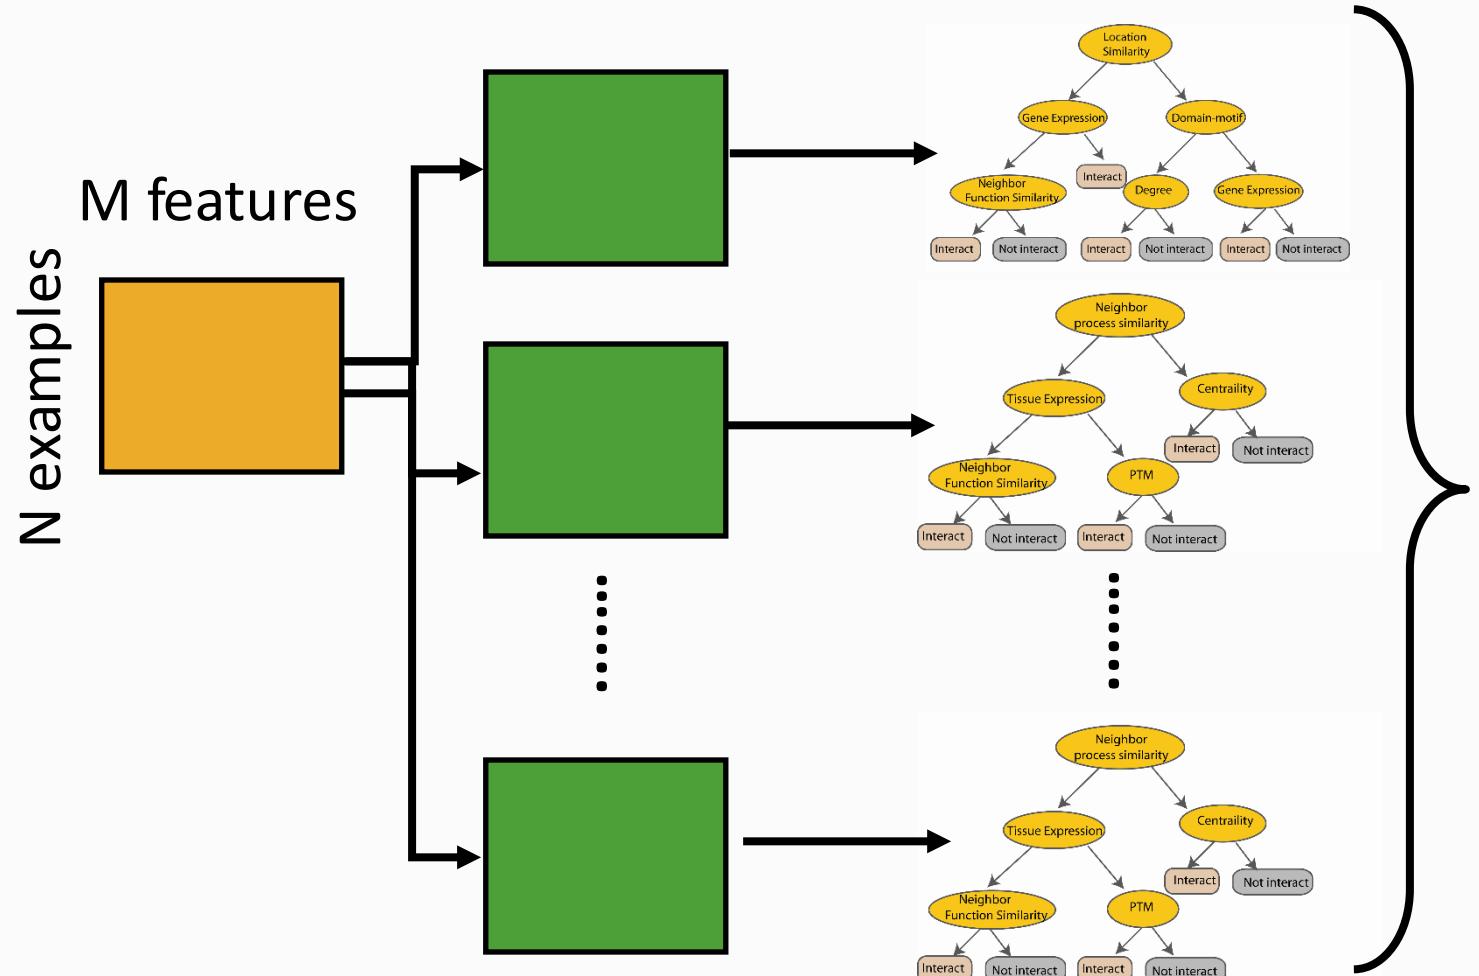
Random Forest Classifier



Random Forest Classifier



Random Forest Classifier



Take the majority vote

Random Forests

- Combination of decision trees and bootstrapping concepts
- A large number of decision trees is trained, each on a different bootstrap sample
- At each split, only a random number of the original variables is available (i.e. small selection of columns)
- Data points are classified by majority voting of the individual trees



Key Criteria for Ensembling

Consensus

- Each person has access to the same information on which to learn and base their decision.

Independence

- People's opinions are not determined by the opinions of those around them.

Decentralization

- People are able to specialize and draw on local knowledge.

Aggregation

- Some mechanism exists to turn private judgments into a collective decision



Key Criteria for Ensembling

Diversity of Opinion

- Each person should have **private information** even if it's just an eccentric interpretation of the known facts.

Independence

- People's opinions are not determined by the opinions of those around them.

Decentralization

- People are able to specialize and draw on local knowledge.

Aggregation

- Some mechanism exists to turn private judgments into a collective decision



So, how good are ensemble methods...

Netflix Prize

Began October 2006

Supervised learning task

- Training data is a set of users and ratings (1,2,3,4,5 stars) those users have given to movies.
- Construct a classifier that given a user and an unrated movie, correctly classifies that movie as either 1, 2, 3, 4, or 5 stars

\$1 million prize for a 10% improvement over Netflix's current movie recommender/classifier
(MSE = 0.9514)

<http://www.wired.com/business/2009/09/how-the-netflix-prize-was-won/>, a light read (highly suggested)



Netflix Prize

Home | Rules | Leaderboard | Register | Update | Submit | Download

Leaderboard

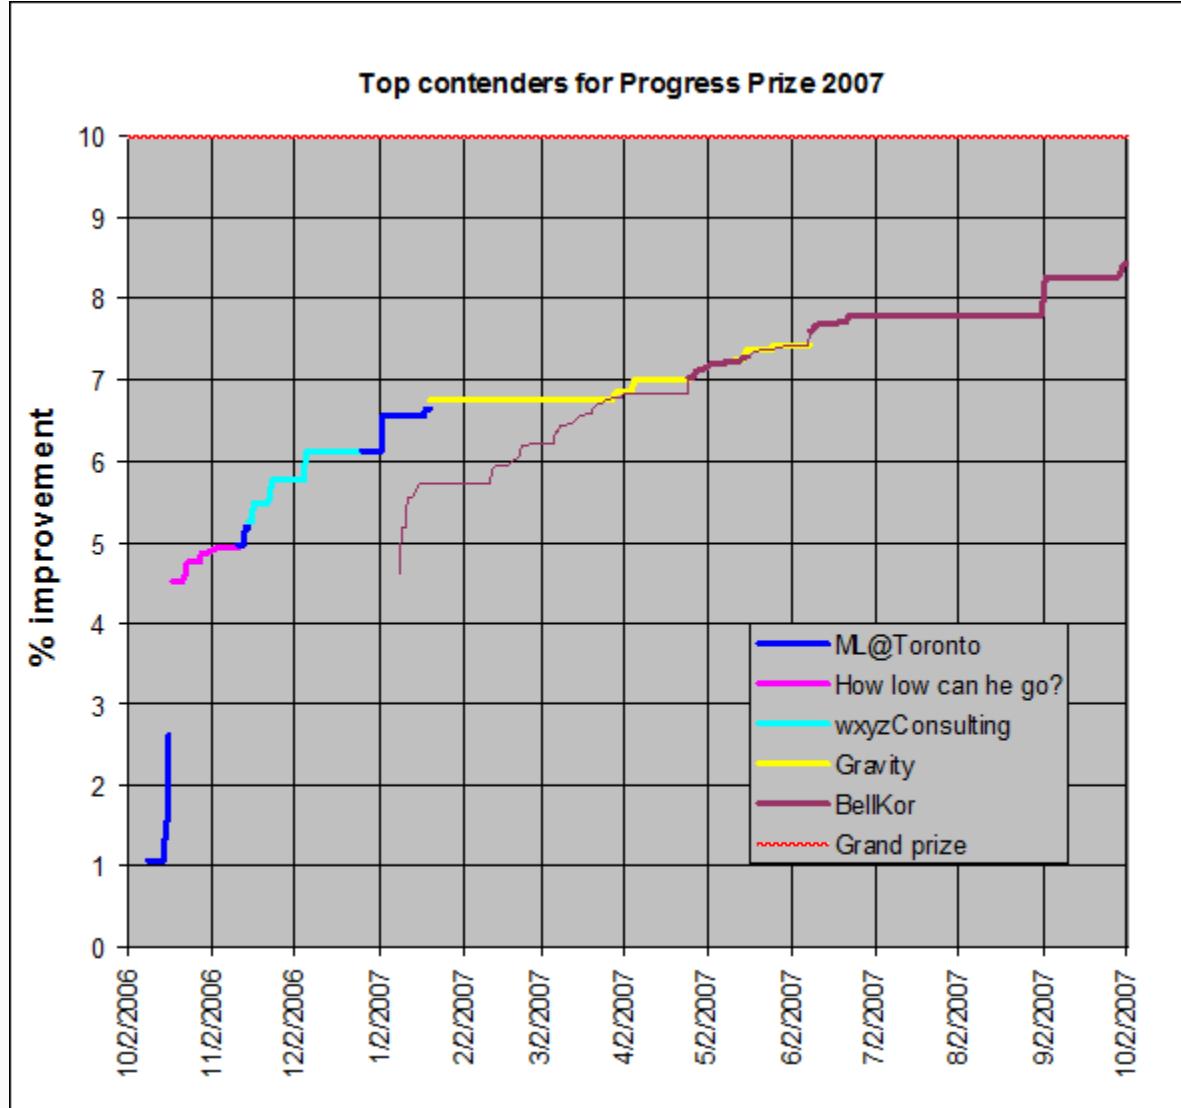
Team Name	Best Score	% Improvement
No Grand Prize candidates yet	-	-
Grand Prize - RMSE <= 0.8563		
How low can he go?	0.9046	4.92
ML@UToronto A	0.9046	4.92
ssorkin	0.9089	4.47
wxyzconsulting.com	0.9103	4.32
The Thought Gang	0.9113	4.21
NIPS Reject	0.9118	4.16
simonfunk	0.9145	3.88
Bozo_The_Clown	0.9177	3.54
Elliptic Chaos	0.9179	3.52
datcracker	0.9183	3.48
Foreseer	0.9214	3.15
bsdfish	0.9229	3.00
Three Blind Mice	0.9234	2.94
Bocsimacko	0.9238	2.90
Remco	0.9252	2.75
karmatics	0.9301	2.24
Chapelator	0.9314	2.10
Flmod	0.9325	1.99
mthrox	0.9328	1.96

Just three weeks after it began, at least 40 teams had bested the Netflix classifier.

Top teams showed about 5% improvement.



However, improvement slowed...



from <http://www.research.att.com/~volinsky/netflix/>

--	No Progress Prize candidates yet	--	--
Progress Prize - RMSE <= 0.8625			
1	BellKor	0.8705	8.50
Progress Prize 2007 - RMSE = 0.8712 - Winning Team: KorBell			
2	KorBell	0.8712	8.43
3	When Gravity and Dinosaurs Unite	0.8717	8.38
4	Gravity	0.8743	8.10
5	basho	0.8746	8.07
6	Dinosaur Planet	0.8753	8.00
7	ML@UToronto A	0.8787	7.64
8	Arek Paterek	0.8789	7.62
9	NIPS Reject	0.8808	7.42
10	Just a guy in a garage	0.8834	7.15
11	Ensemble Experts	0.8841	7.07
12	mathematical capital	0.8844	7.04
13	HowLowCanHeGo2	0.8847	7.01
14	The Thought Gang	0.8849	6.99
15	Reel Ingenuity	0.8855	6.93
16	strudeltamale	0.8859	6.88
17	NIPS Submission	0.8861	6.86
18	Three Blind Mice	0.8869	6.78
19	TrainOnTest	0.8869	6.78
20	Geoff Dean	0.8869	6.78
21	Rookies	0.8872	6.75
22	Paul Harrison	0.8872	6.75
23	ATTEAM	0.8873	6.74
24	wxyzconsulting.com	0.8874	6.73
25	ICMLsubmission	0.8875	6.72
26	Efratko	0.8877	6.70
27	Kitty	0.8881	6.65
28	SecondaryResults	0.8884	6.62
29	Birgit Kraft	0.8885	6.61

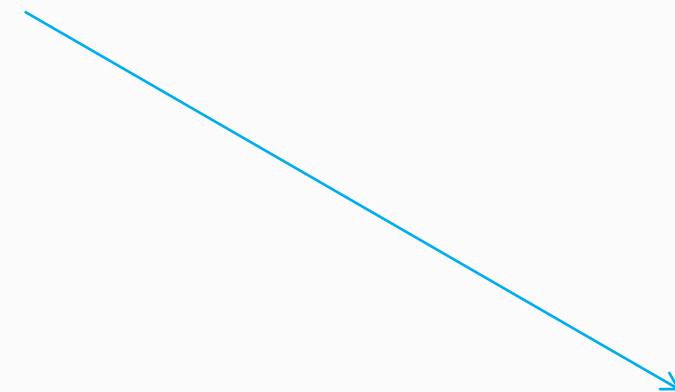
The top team posted a 8.5% improvement.

Ensemble methods are the best performers...



Rookies

“Thanks to Paul Harrison's collaboration, a simple mix of our solutions improved our result from 6.31 to 6.75”



No Progress Prize candidates yet			
Progress Prize - RMSE <= 0.8625			
1	BellKor	0.8705	8.50
Progress Prize 2007 - RMSE = 0.8712 - Winning Team: KorBell			
2	KorBell	0.8712	8.43
3	When Gravity and Dinosaurs Unite	0.8717	8.38
4	Gravity	0.8743	8.10
5	basho	0.8746	8.07
6	Dinosaur Planet	0.8753	8.00
7	ML@UToronto A	0.8787	7.64
8	Arek Paterek	0.8789	7.62
9	NIPS Reject	0.8808	7.42
10	Just a guy in a garage	0.8834	7.15
11	Ensemble Experts	0.8841	7.07
12	mathematical capital	0.8844	7.04
13	HowLowCanHeGo2	0.8847	7.01
14	The Thought Gang	0.8849	6.99
15	Reel Ingenuity	0.8855	6.93
16	strudeltamale	0.8859	6.88
17	NIPS Submission	0.8861	6.86
18	Three Blind Mice	0.8869	6.78
19	TrainOnTest	0.8869	6.78
20	Geoff.Dean	0.8869	6.78
21	Rookies	0.8872	6.75
22	Paul Harrison	0.8872	6.75
23	ATTEAM	0.8873	6.74
24	wxyzconsulting.com	0.8874	6.73
25	ICMLsubmission	0.8875	6.72
26	Efratko	0.8877	6.70
27	Kitty	0.8881	6.65
28	SecondaryResults	0.8884	6.62
29	Birgit Kraft	0.8885	6.61

Arek Paterek

“My approach is to **combine the results of many methods** (also two-way interactions between them) using linear regression on the test set. The best method in my ensemble is regularized SVD with biases, post processed with kernel ridge regression”

http://rainbow.mimuw.edu.pl/~ap/ap_kdd.pdf

No Progress Prize candidates yet			
Progress Prize - RMSE <= 0.8625			
1	BellKor	0.8705	8.50
Progress Prize 2007 - RMSE = 0.8712 - Winning Team: KorBell			
2	KorBell	0.8712	8.43
3	When Gravity and Dinosaurs Unite	0.8717	8.38
4	Gravity	0.8743	8.10
5	basho	0.8746	8.07
6	Dinosaur Planet	0.8753	8.00
7	ML@UToronto A	0.8787	7.64
8	Arek Paterek	0.8789	7.62
9	NIPS Reject	0.8808	7.42
10	Just a guy in a garage	0.8834	7.15
11	Ensemble Experts	0.8841	7.07
12	mathematical capital	0.8844	7.04
13	HowLowCanHeGo2	0.8847	7.01
14	The Thought Gang	0.8849	6.99
15	Reel Ingenuity	0.8855	6.93
16	strudeltamale	0.8859	6.88
17	NIPS Submission	0.8861	6.86
18	Three Blind Mice	0.8869	6.78
19	TrainOnTest	0.8869	6.78
20	Geoff Dean	0.8869	6.78
21	Rookies	0.8872	6.75
22	Paul Harrison	0.8872	6.75
23	ATTEAM	0.8873	6.74
24	wxyzconsulting.com	0.8874	6.73
25	ICMLsubmission	0.8875	6.72
26	Efratko	0.8877	6.70
27	Kitty	0.8881	6.65
28	SecondaryResults	0.8884	6.62
29	Birgit Kraft	0.8885	6.61

U of Toronto

“When the predictions of **multiple** RBM models and **multiple** SVD models are linearly combined, we achieve an error rate that is well over 6% better than the score of Netflix’s own system.”

<http://www.cs.toronto.edu/~rsalakhu/papers/rbmcf.pdf>

No Progress Prize candidates yet			
Progress Prize - RMSE <= 0.8625			
1	BellKor	0.8705	8.50
Progress Prize 2007 - RMSE = 0.8712 - Winning Team: KorBell			
2	KorBell	0.8712	8.43
3	When Gravity and Dinosaurs Unite	0.8717	8.38
4	Gravity	0.8743	8.10
5	basho	0.8746	8.07
6	Dinosaur Planet	0.8753	8.00
7	ML@UToronto A	0.8787	7.64
8	Arek Paterek	0.8789	7.62
9	NIPS Reject	0.8808	7.42
10	Just a guy in a garage	0.8834	7.15
11	Ensemble Experts	0.8841	7.07
12	mathematical capital	0.8844	7.04
13	HowLowCanHeGo2	0.8847	7.01
14	The Thought Gang	0.8849	6.99
15	Reel Ingenuity	0.8855	6.93
16	strudeltamale	0.8859	6.88
17	NIPS Submission	0.8861	6.86
18	Three Blind Mice	0.8869	6.78
19	TrainOnTest	0.8869	6.78
20	Geoff Dean	0.8869	6.78
21	Rookies	0.8872	6.75
22	Paul Harrison	0.8872	6.75
23	ATTEAM	0.8873	6.74
24	wxyzconsulting.com	0.8874	6.73
25	ICMLsubmission	0.8875	6.72
26	Efratko	0.8877	6.70
27	Kitty	0.8881	6.65
28	SecondaryResults	0.8884	6.62
29	Birgit Kraft	0.8885	6.61

Gravity

Table 5: Best results of single approaches and their combinations

Method/Combination	RMSE
MF	0.9190
NB	0.9313
CL	0.9606
NB + CL	0.9275
MF + CL	0.9137
MF + NB	0.9089
MF + NB + CL	0.9089

home.mit.bme.hu/~gtakacs/download/gravity.pdf



No Progress Prize candidates yet			
Progress Prize - RMSE <= 0.8625			
1	BellKor	0.8705	8.50
Progress Prize 2007 - RMSE = 0.8712 - Winning Team: KorBell			
2	KorBell	0.8712	8.43
3	When Gravity and Dinosaurs Unite	0.8717	8.38
4	Gravity	0.8743	8.10
5	basho	0.8746	8.07
6	Dinosaur Planet	0.8753	8.00
7	ML@UToronto A	0.8787	7.64
8	Arek Paterek	0.8789	7.62
9	NIPS Reject	0.8808	7.42
10	Just a guy in a garage	0.8834	7.15
11	Ensemble Experts	0.8841	7.07
12	mathematical capital	0.8844	7.04
13	HowLowCanHeGo2	0.8847	7.01
14	The Thought Gang	0.8849	6.99
15	Reel Ingenuity	0.8855	6.93
16	strudeltamale	0.8859	6.88
17	NIPS Submission	0.8861	6.86
18	Three Blind Mice	0.8869	6.78
19	TrainOnTest	0.8869	6.78
20	Geoff Dean	0.8869	6.78
21	Rookies	0.8872	6.75
22	Paul Harrison	0.8872	6.75
23	ATTEAM	0.8873	6.74
24	wxyzconsulting.com	0.8874	6.73
25	ICMLsubmission	0.8875	6.72
26	Efratko	0.8877	6.70
27	Kitty	0.8881	6.65
28	SecondaryResults	0.8884	6.62
29	Birgit Kraft	0.8885	6.61

When Gravity and Dinosaurs Unite

“Our common team blends the result of team Gravity and team Dinosaur Planet.”

Might have guessed from the name...

No Progress Prize candidates yet			
Progress Prize - RMSE <= 0.8625			
1	BellKor	0.8705	8.50
Progress Prize 2007 - RMSE = 0.8712 - Winning Team: KorBell			
2	KorBell	0.8712	8.43
3	When Gravity and Dinosaurs Unite	0.8717	8.38
4	Gravity	0.8743	8.10
5	basho	0.8746	8.07
6	Dinosaur Planet	0.8753	8.00
7	ML@UToronto A	0.8787	7.64
8	Arek Paterek	0.8789	7.62
9	NIPS Reject	0.8808	7.42
10	Just a guy in a garage	0.8834	7.15
11	Ensemble Experts	0.8841	7.07
12	mathematical capital	0.8844	7.04
13	HowLowCanHeGo2	0.8847	7.01
14	The Thought Gang	0.8849	6.99
15	Reel Ingenuity	0.8855	6.93
16	strudeltamale	0.8859	6.88
17	NIPS Submission	0.8861	6.86
18	Three Blind Mice	0.8869	6.78
19	TrainOnTest	0.8869	6.78
20	Geoff Dean	0.8869	6.78
21	Rookies	0.8872	6.75
22	Paul Harrison	0.8872	6.75
23	ATTEAM	0.8873	6.74
24	wxyzconsulting.com	0.8874	6.73
25	ICMLsubmission	0.8875	6.72
26	Efratko	0.8877	6.70
27	Kitty	0.8881	6.65
28	SecondaryResults	0.8884	6.62
29	Birgit Kraft	0.8885	6.61

BellKor / KorBell

And, yes, the top team which is from AT&T...

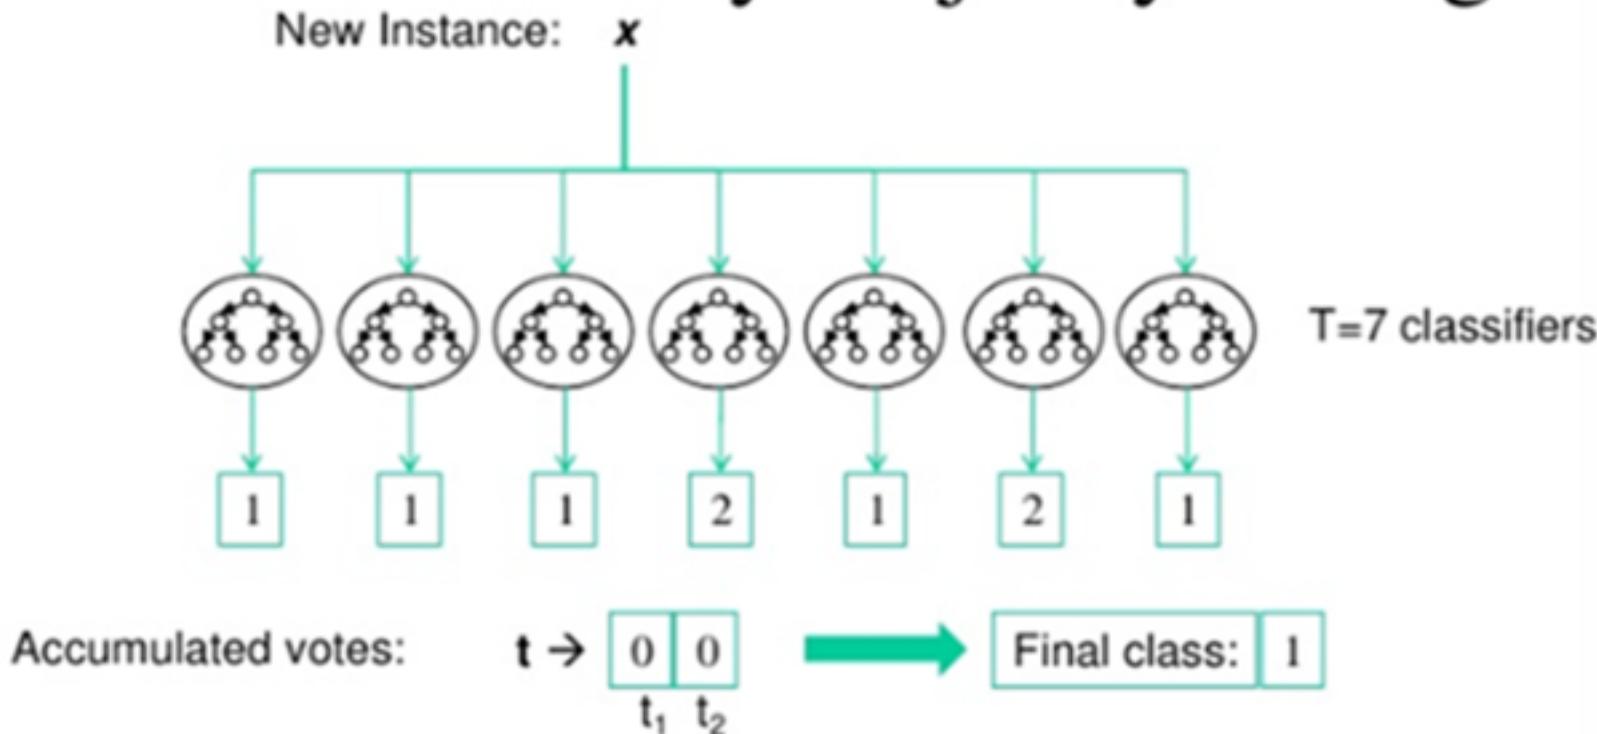
“Our final solution (RMSE=0.8712) consists
of blending 107 individual results.”

No Progress Prize candidates yet			
Progress Prize - RMSE <= 0.8625			
1	BellKor	0.8705	8.50
Progress Prize 2007 - RMSE = 0.8712 - Winning Team: KorBell			
2	KorBell	0.8712	8.43
3	When Gravity and Dinosaurs Unite	0.8717	8.38
4	Gravity	0.8743	8.10
5	basho	0.8746	8.07
6	Dinosaur Planet	0.8753	8.00
7	ML@UToronto A	0.8787	7.64
8	Arek Paterek	0.8789	7.62
9	NIPS Reject	0.8808	7.42
10	Just a guy in a garage	0.8834	7.15
11	Ensemble Experts	0.8841	7.07
12	mathematical capital	0.8844	7.04
13	HowLowCanHeGo2	0.8847	7.01
14	The Thought Gang	0.8849	6.99
15	Reel Ingenuity	0.8855	6.93
16	strudeltamale	0.8859	6.88
17	NIPS Submission	0.8861	6.86
18	Three Blind Mice	0.8869	6.78
19	TrainOnTest	0.8869	6.78
20	Geoff Dean	0.8869	6.78
21	Rookies	0.8872	6.75
22	Paul Harrison	0.8872	6.75
23	ATTEAM	0.8873	6.74
24	wxyzconsulting.com	0.8874	6.73
25	ICMLsubmission	0.8875	6.72
26	Efratko	0.8877	6.70
27	Kitty	0.8881	6.65
28	SecondaryResults	0.8884	6.62
29	Birgit Kraft	0.8885	6.61

Ensemble Classification

Aggregation of predictions of multiple classifiers with the goal of improving accuracy.

Classification by majority voting



Bagging (Bootstrap Aggregating)

Given

- Training set of N examples
- A class of learning models (e.g. decision trees, neural networks, ...)

Method

- Train multiple (k) models on different samples (data splits) and average their predictions
- Predict (test) by averaging the results of k models

Goal

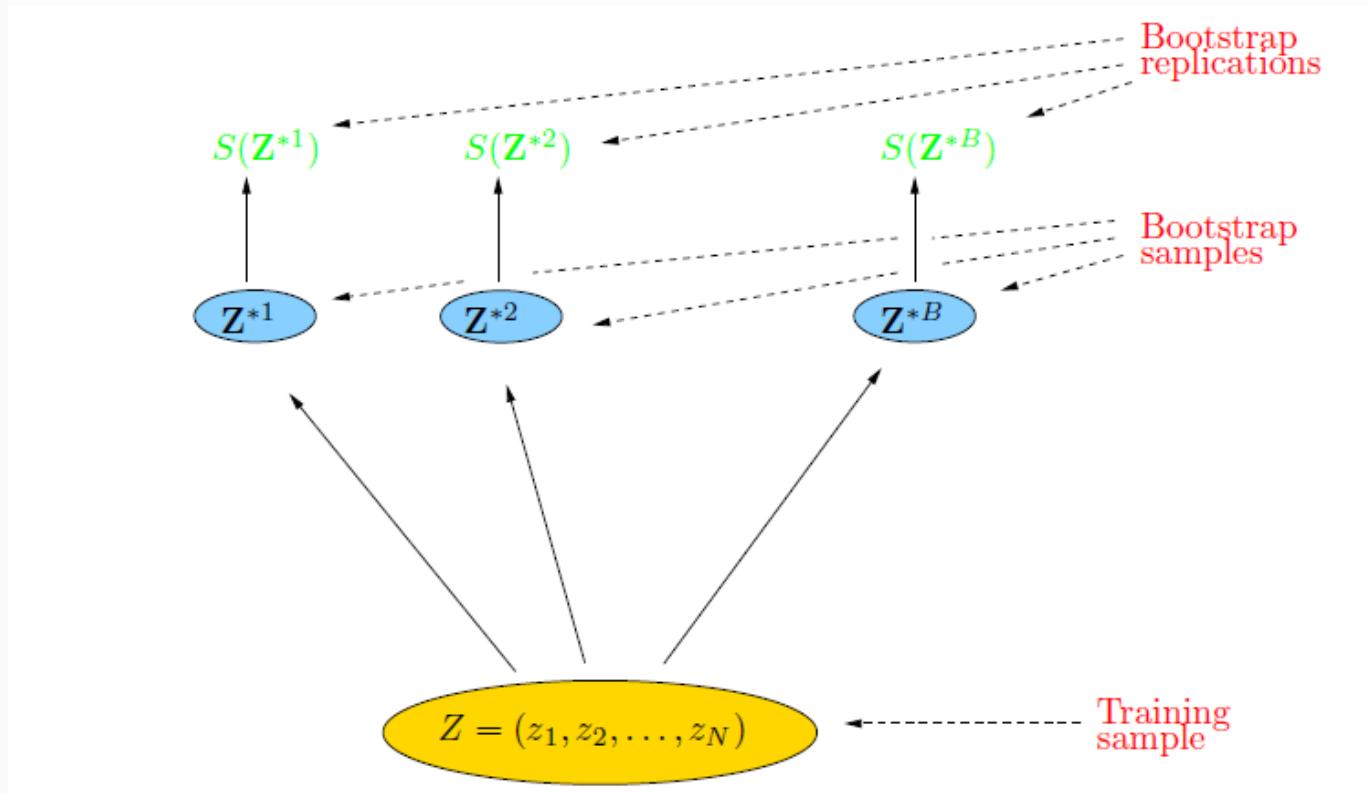
- Improve the accuracy of one model by using its multiple copies
- Average of misclassification errors on different data split gives a better estimate of the predictive ability of a learning method



Bootstrap

The basic idea:

Randomly draw datasets *with replacement* from the training data, each sample *the same size as the original training set*



Bagging Algorithm

Training

- In each iteration t , $t=1,\dots,T$
 - Randomly sample with replacement N samples from the training set
 - Train a chosen “base model” (e.g. neural network, decision tree) on the samples

Test

- For each test example
 - Start all trained base models
 - Predict by combining results of all T trained models:
 - **Regression**: averaging
 - **Classification**: a majority vote

More on Bagging

Bagging works because it reduces variance by voting/averaging

- Note: in some hypothetical situations the overall error might increase;
- Usually, the more classifiers the better;

Can help a lot if data is noisy

Can also be applied to numeric prediction

If the learning algorithm is unstable, then bagging almost always improves performance. Bagging stable classifiers is not a good idea

- Which ones are unstable? Neural nets, decision trees, regression trees, linear regression
- Which ones are stable? K-nearest neighbors



What is Boosting?

- Analogy: Consult several doctors, based on a combination of weighted diagnoses—weight assigned based on the previous diagnosis accuracy
- How boosting works?
 - Weights are assigned to each training tuple
 - A series of k classifiers is iteratively learned
 - After a classifier M_i is learned, the weights are updated to allow the subsequent classifier, M_{i+1} , to pay more attention to the training tuples that were misclassified by M_i
 - The final M^* combines the votes of each individual classifier, where the weight of each classifier's vote is a function of its accuracy
- Boosting algorithm can be extended for the prediction of continuous values
- Comparing with bagging: boosting tends to achieve greater accuracy, but it also risks overfitting the model to misclassified data



The Basic Idea

- Suppose there are just 5 training examples $\{1,2,3,4,5\}$
- Initially each example has a 0.2 ($1/5$) probability of being sampled
- 1st round of boosting samples (with replacement) 5 examples:
- $\{2, 4, 4, 3, 2\}$ and builds a classifier from them
- Suppose examples 2, 3, 5 are correctly predicted by this classifier, and examples 1, 4 are wrongly predicted:
 - Weight of examples 1 and 4 is increased,
Weight of examples 2, 3, 5 is decreased
- 2nd round of boosting samples again 5 examples, but now examples 1 and 4 are more likely to be sampled
- And so on ...until some convergence is achieved



Boosting

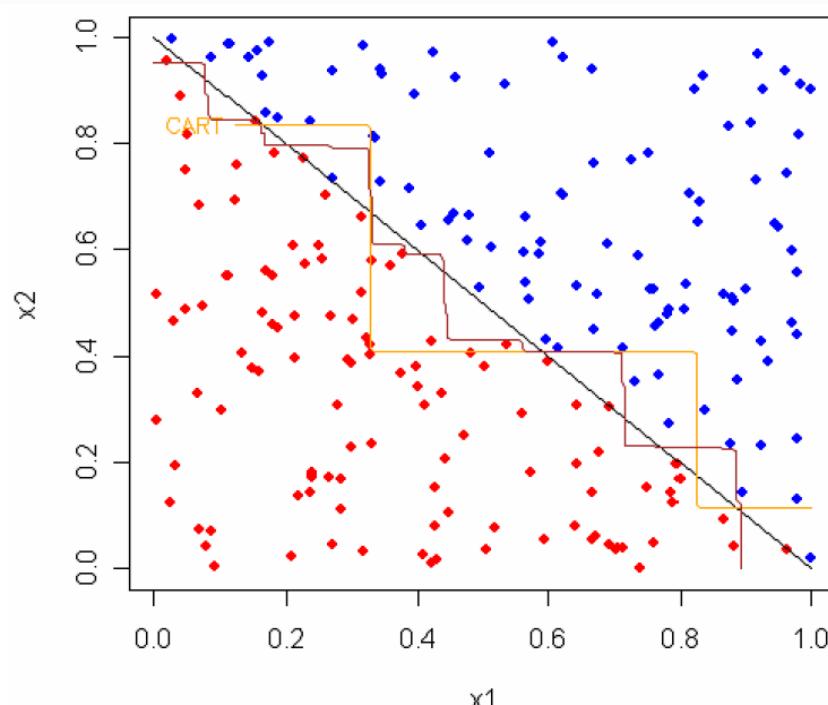
- Also uses voting/averaging
- Weights models according to performance
- Iterative: new models are influenced by the performance of previously built ones
 - Encourage new model to become an “expert” for instances misclassified by earlier models
 - Intuitive justification: models should be experts that complement each other
- Several variants
 - Boosting by sampling, the weights are used to sample the data for training
 - Boosting by weighting, the weights are used by the learning algorithm



Bagging: reduces variance – Example 1

- Two categories of samples: blue, red
- Two predictors: x_1 and x_2

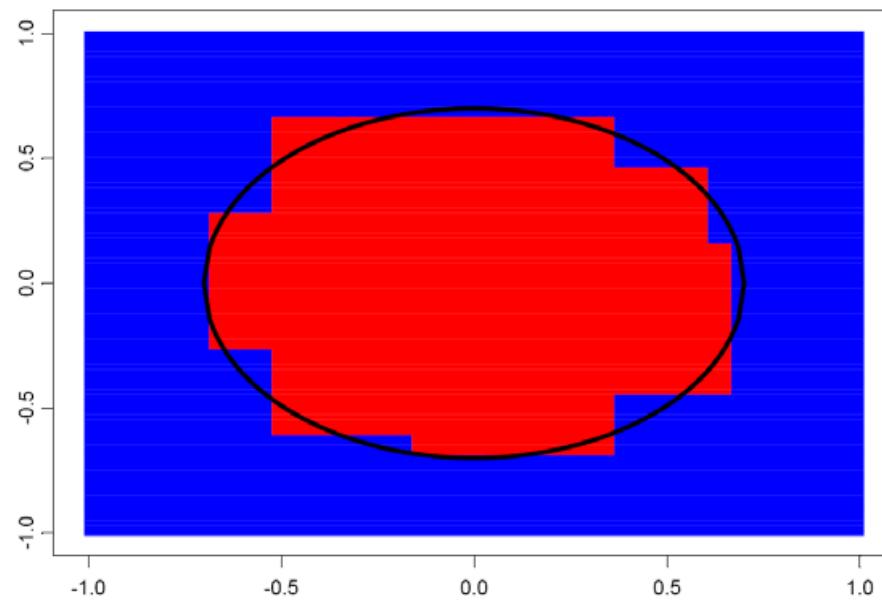
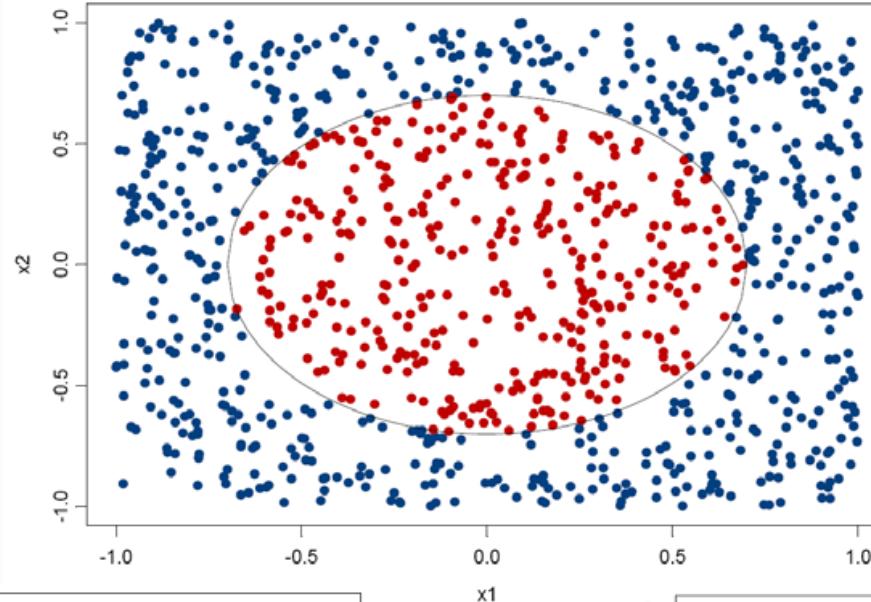
Diagonal separation...hardest case for tree-based classifier



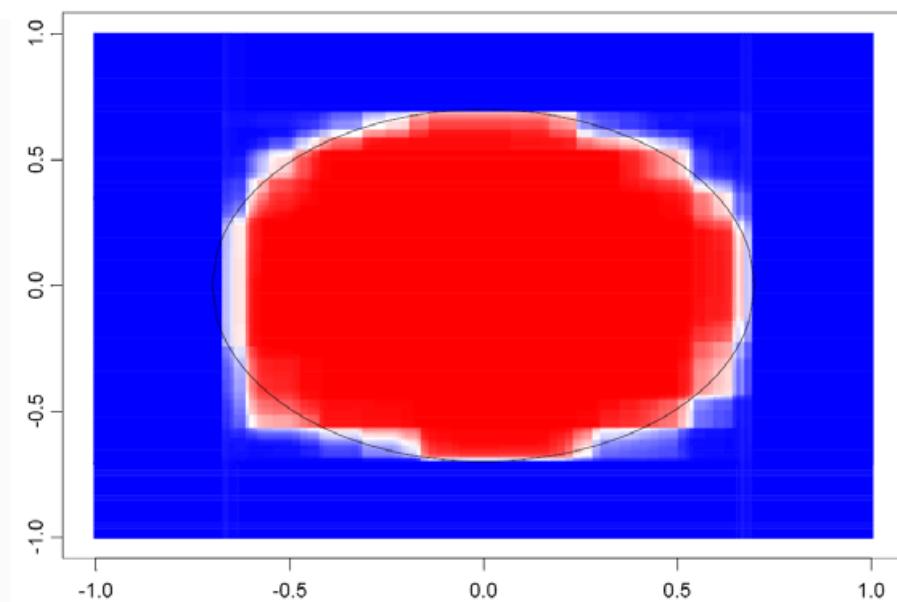
- Single tree decision boundary in orange.
- Bagged predictor decision boundary in red.

Bagging: reduces variance – Example 2

Ellipsoid separation →
Two categories,
Two predictors



Single tree decision boundary

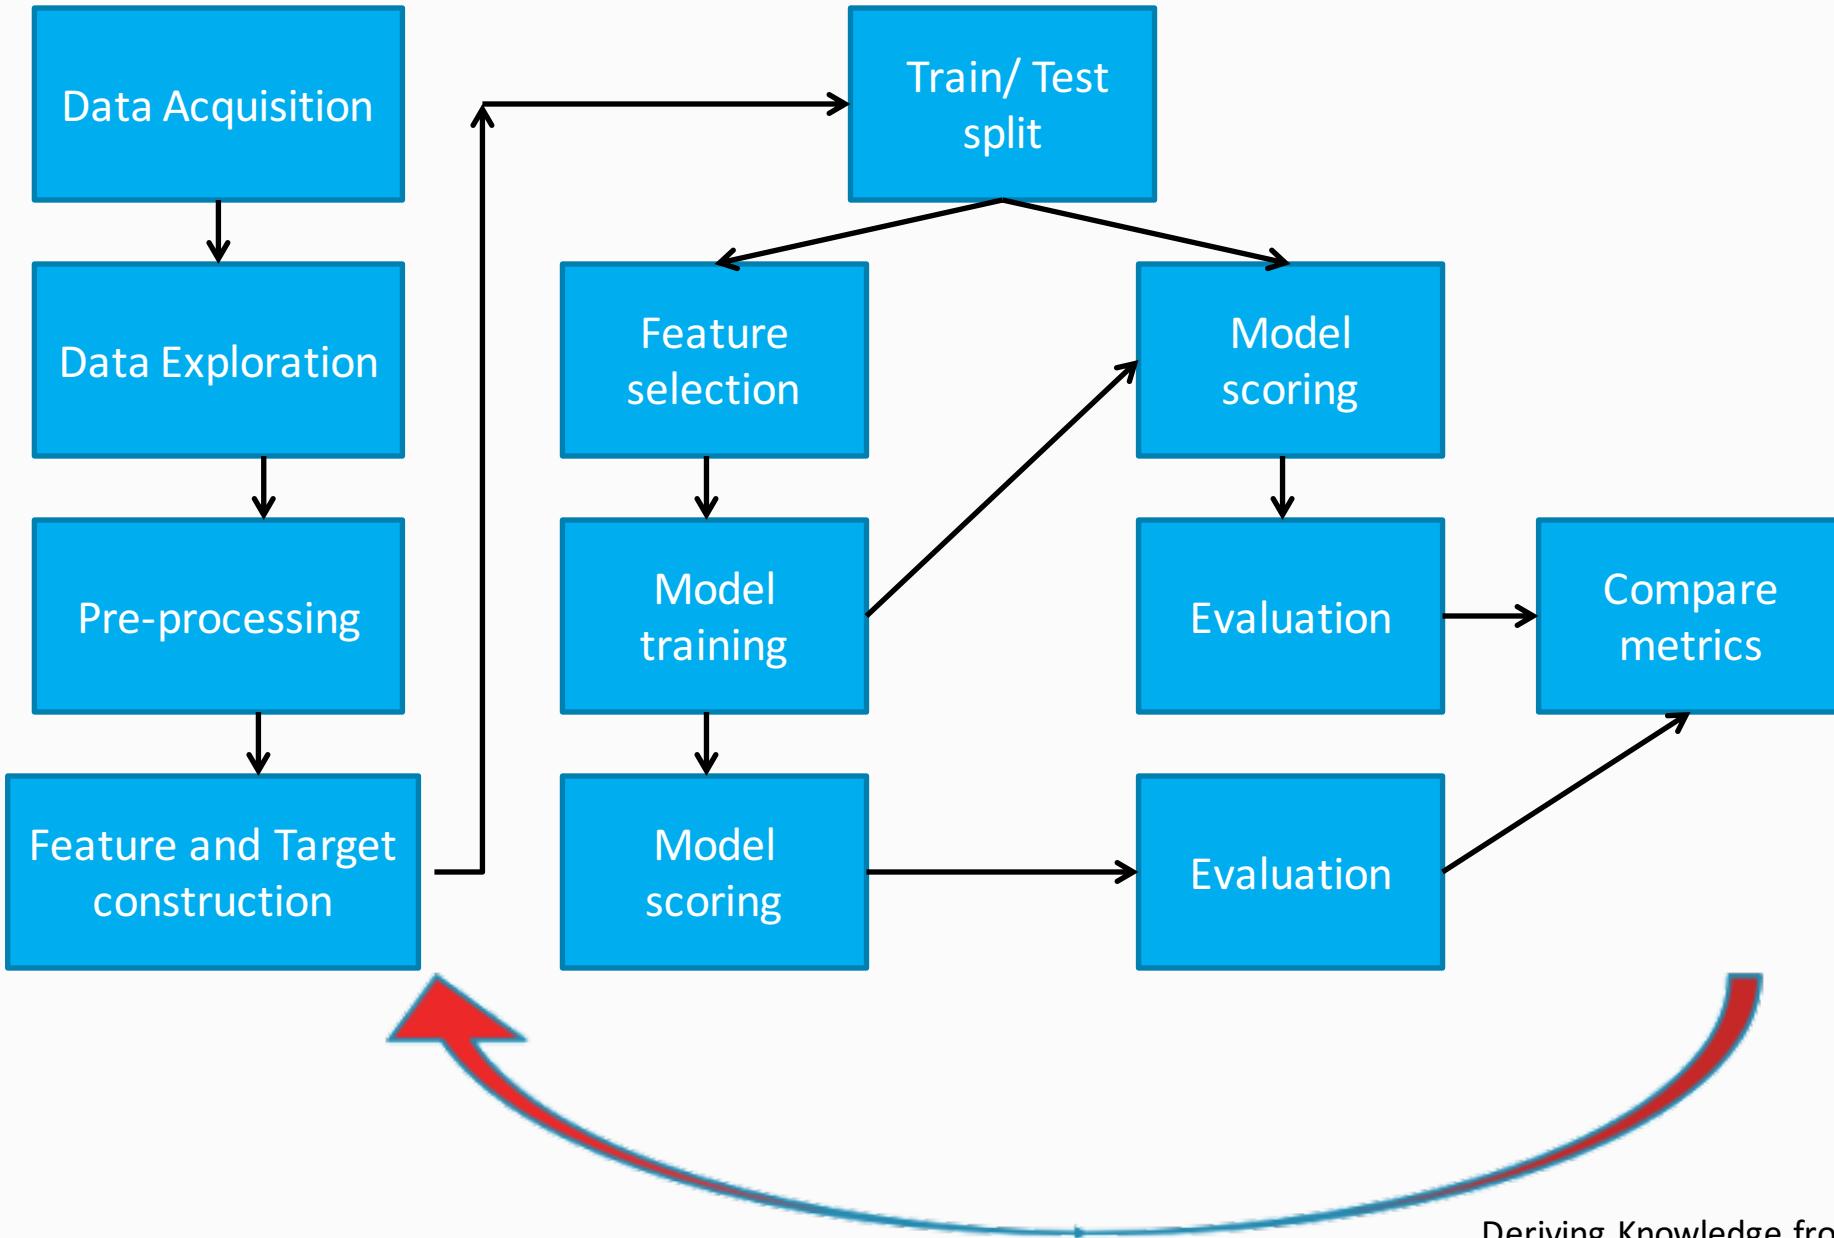


100 bagged trees...Deriving Knowledge from Data at Scale

Lecture 4 Sneak Preview - Model Evaluation



Steps in a Predictive Modeling process



Model Scoring (subject for today...)

Process of applying the model parameters to a dataset of features to generate predictions

Some algorithms produce calibrated scores

- i.e. scores represent probability of data point belonging to a class
e.g. logistic regression, versions of decision trees, Naïve Bayes

Other algorithms produce rank ordering

- i.e. scores can be used for relative ordering of records e.g. SVM



Model Evaluation

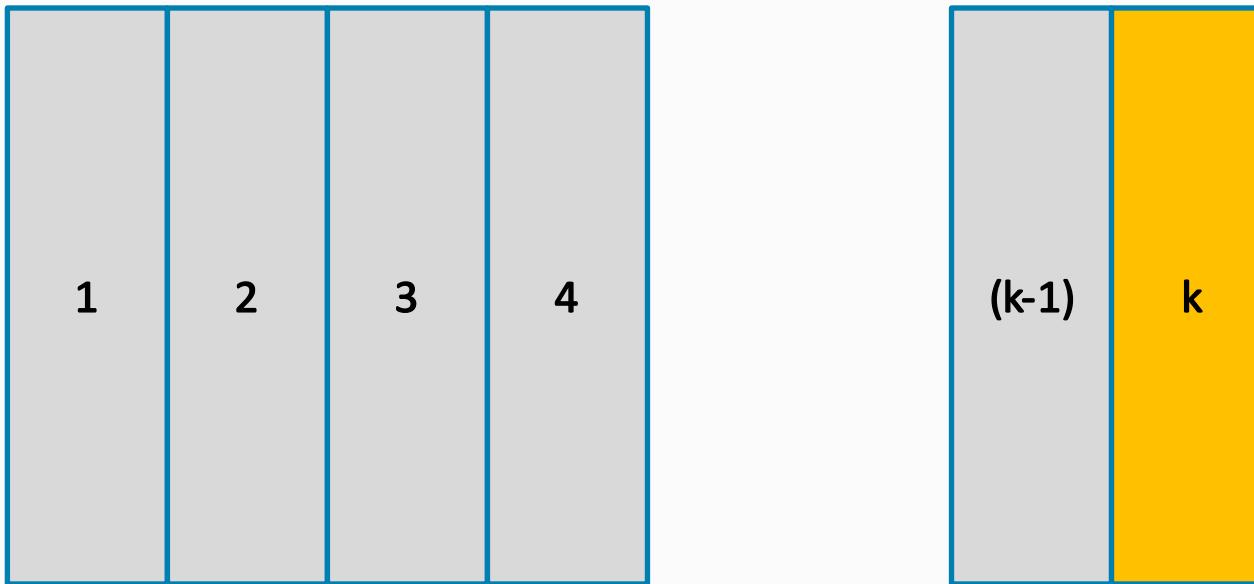
- Compute model performance metrics based on the scores and the true target label
- Metrics measure ability of model to learn the relationship between features and targets



Cross Validation

- Technique for assessing generalization capability of a model, i.e.
 - How well will the model perform on new (unseen) data (drawn from the same distribution as the training data)
- Basic idea is to split the training data into “k” independent pieces (called folds)
 - Train on $(k-1)$ folds and test on the remaining fold
 - Repeat this “k” times, testing once on each fold
 - Average the model and performance metrics from each of these “k” runs
- Typically, $k \sim 10$

Cross Validation



Train

Test

Performance Metrics

Classification

- Consider a two (2) class classification problem.
- Algorithms predict either a
 - **Class** of the example
 - Decision Trees assign the dominant class of the node where the example falls.
 - **Score** of a class for the example
 - Higher the score, greater is the probability of the data point being positive
 - A better model segregates the classes better
 - E.g. Logistic regression, Naïve Bayes, SVM



Performance Metrics (Example)

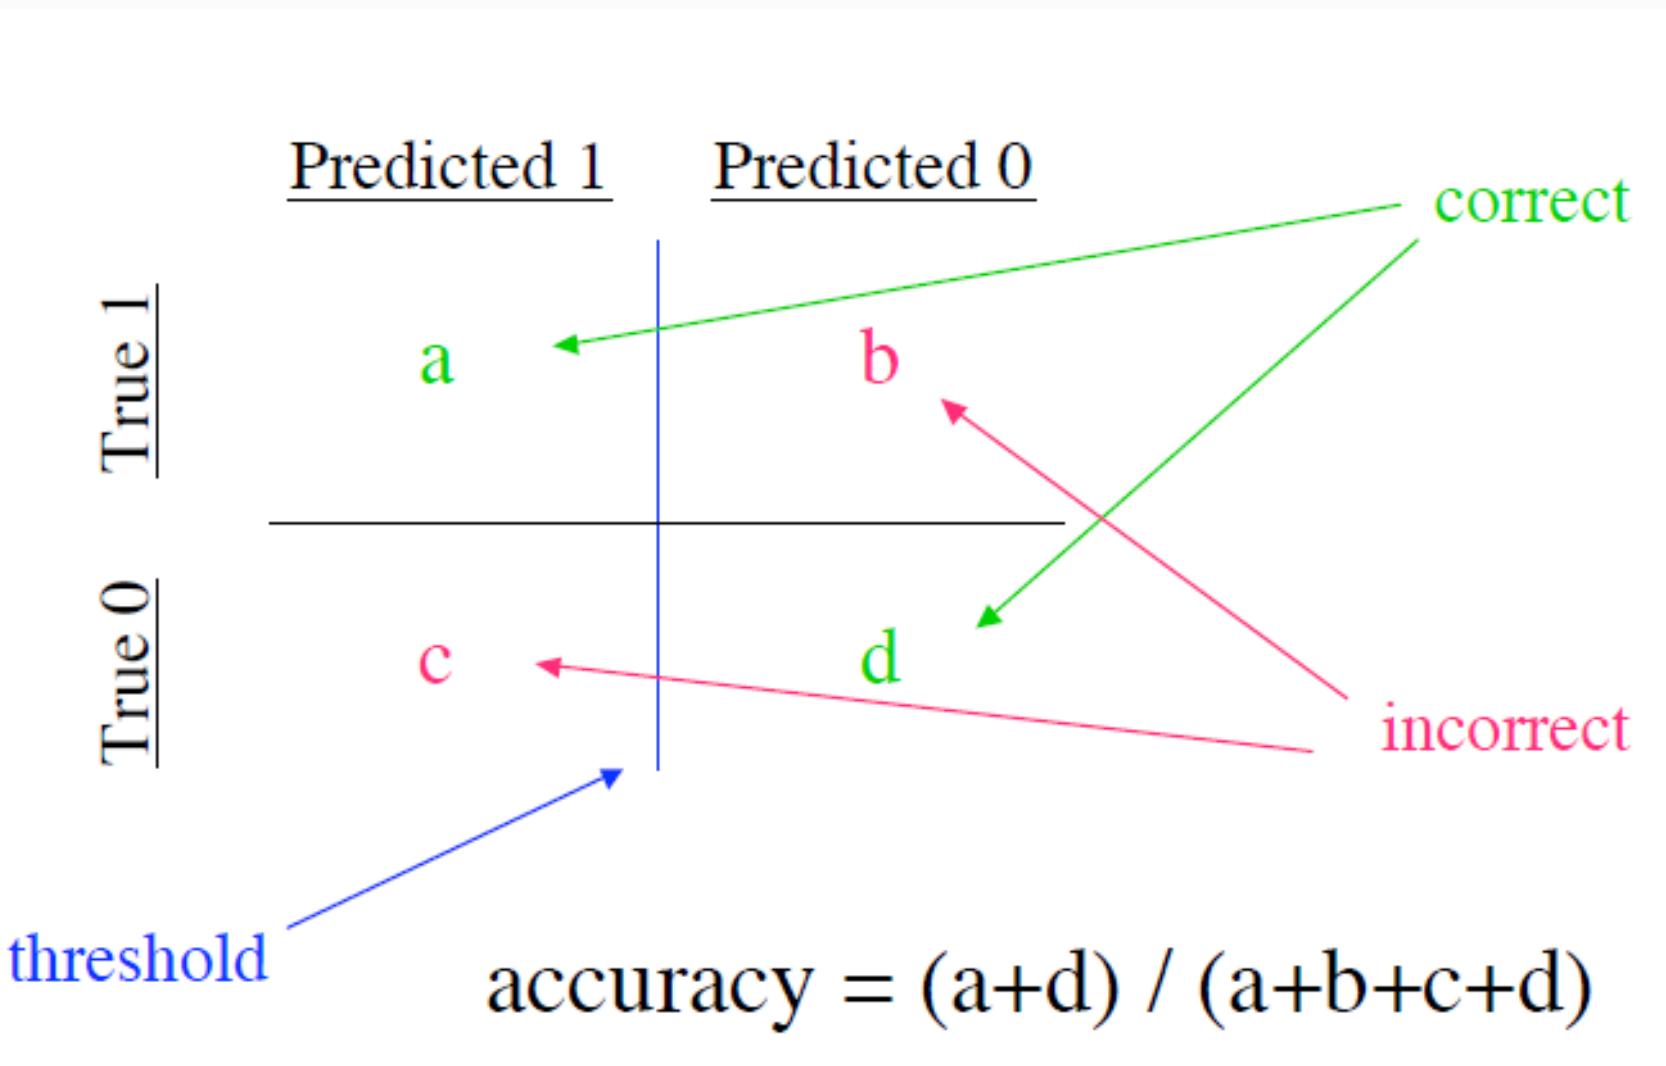
Predicted Label

True Label

Confusion matrix

Row ID	1	0
1	506	122
0	169	420

Performance Metrics



Is Accuracy a good evaluation metric?



Is Accuracy the best evaluation metric?

Accuracy is not the best evaluation metric...

What's wrong with accuracy?

If the vast majority is of binary outcomes are 1, then a stupid model can be accurate but not useful (guess it's always “1”), and a better model might have lower accuracy.



Performance Metrics

Percent Reduction in Error

- 80% accuracy = 20% error
 - *Suppose learning increases accuracy from 80% to 90% error reduced from 20% to 10%*
 - 50% reduction in error



Performance Metrics

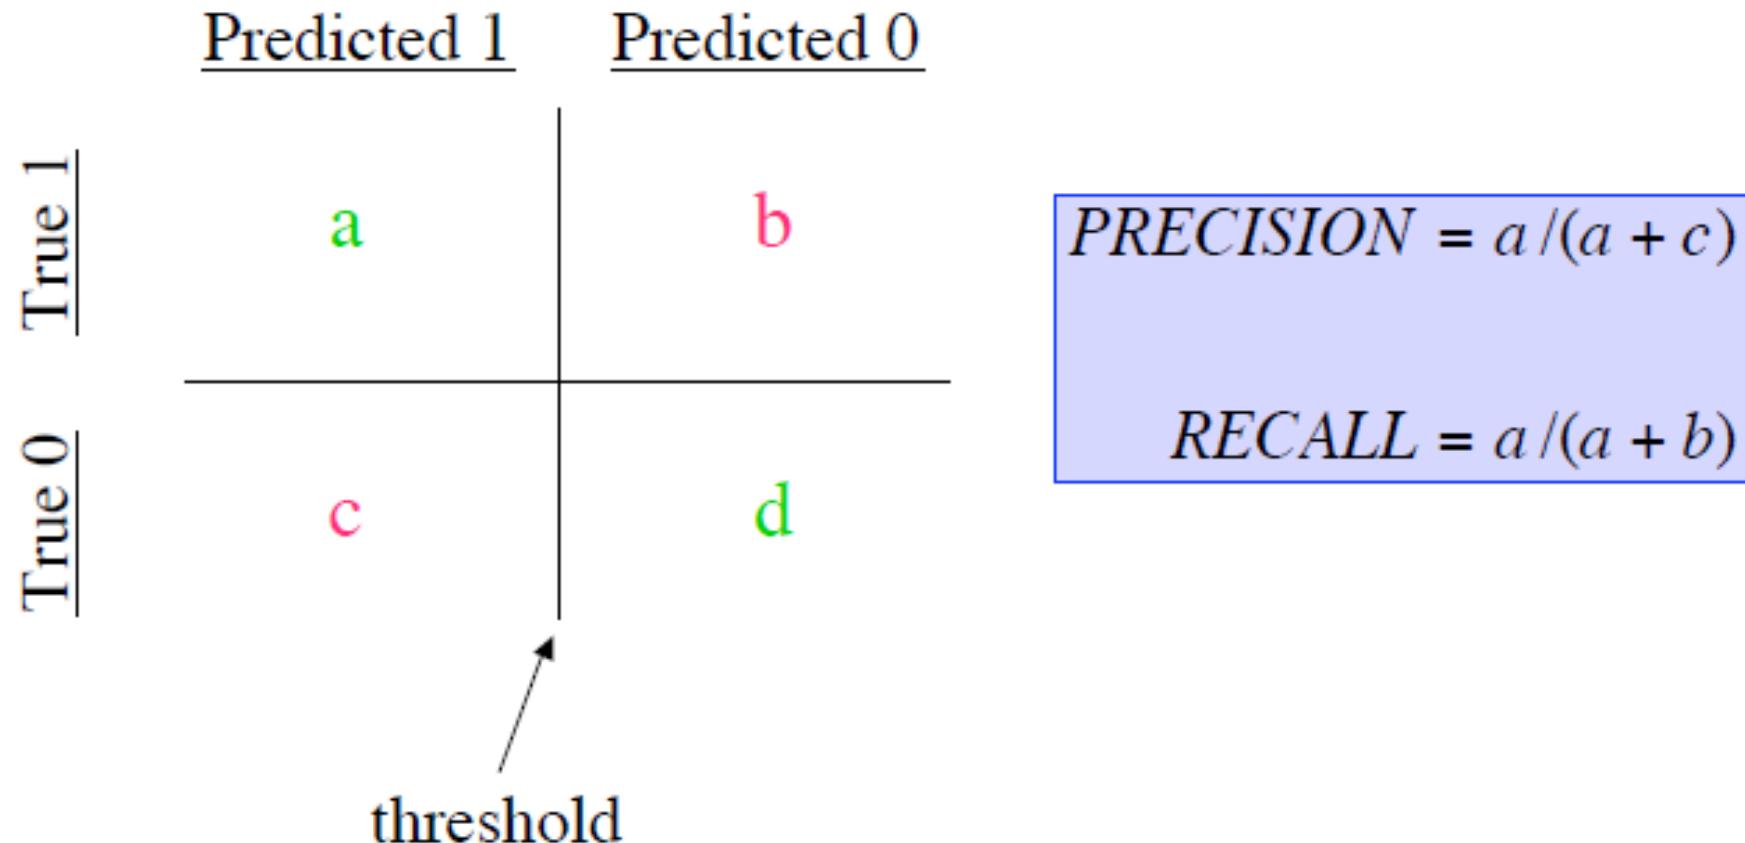
Precision and Recall

- Typically used in document retrieval
- Precision:
 - how many of the returned documents are correct
 - precision (threshold)
- Recall:
 - how many of the positives does the model return
 - recall (threshold)



Performance Metrics

Precision and Recall



$$\text{PRECISION} = a/(a + c)$$

$$\text{RECALL} = a/(a + b)$$

Performance measures in context...

		actual	
		+	-
classifier	+	TP	FP <i>Type I error</i>
	-	FN	TN

		actual	
		+	-
classifier	+	TP	FP
	-	FN	TN

accuracy (ACC)

		actual	
		+	-
classifier	+	TP	FP
	-	FN	TN

neg. predictive value (NPV)

		actual	
		+	-
classifier	+	TP	FP
	-	FN	TN

specificity (SPC)

		actual	
		+	-
classifier	+	TP	FP
	-	FN	TN

true pos rate (TPR)
≡ sensitivity
≡ recall

		actual	
		+	-
classifier	+	TP	FP
	-	FN	TN

false pos rate (FPR)

ROC curve

		actual	
		+	-
classifier	+	TP	FP
	-	FN	TN

pos. predictive value (PPV)
≡ precision

↔ "one minus"

		actual	
		+	-
classifier	+	TP	FP
	-	FN	TN

false discovery rate (FDR)

Dark color is numerator, dark and light color is denominator.

Data Science

Deriving Knowledge from Data at Scale

That's all for tonight...

